

# OpenShift 4

## *Making introductions*

**Samuel Terburg**  
**Red Hat Certified Architect**  
**Cloud-Native Expert**

2020-01-14

A comprehensive overview of “OpenShift Container Platform”, accompanied with interactive Lab exercises.

# Do's & Don'ts

## Do's

### **Manage time**

09:00 Start

12:00 - 13:30 Lunch

16:00 Go Home

### **Ask Questions**

**Fantasize** (real-world cases)

**Have Fun** (share)

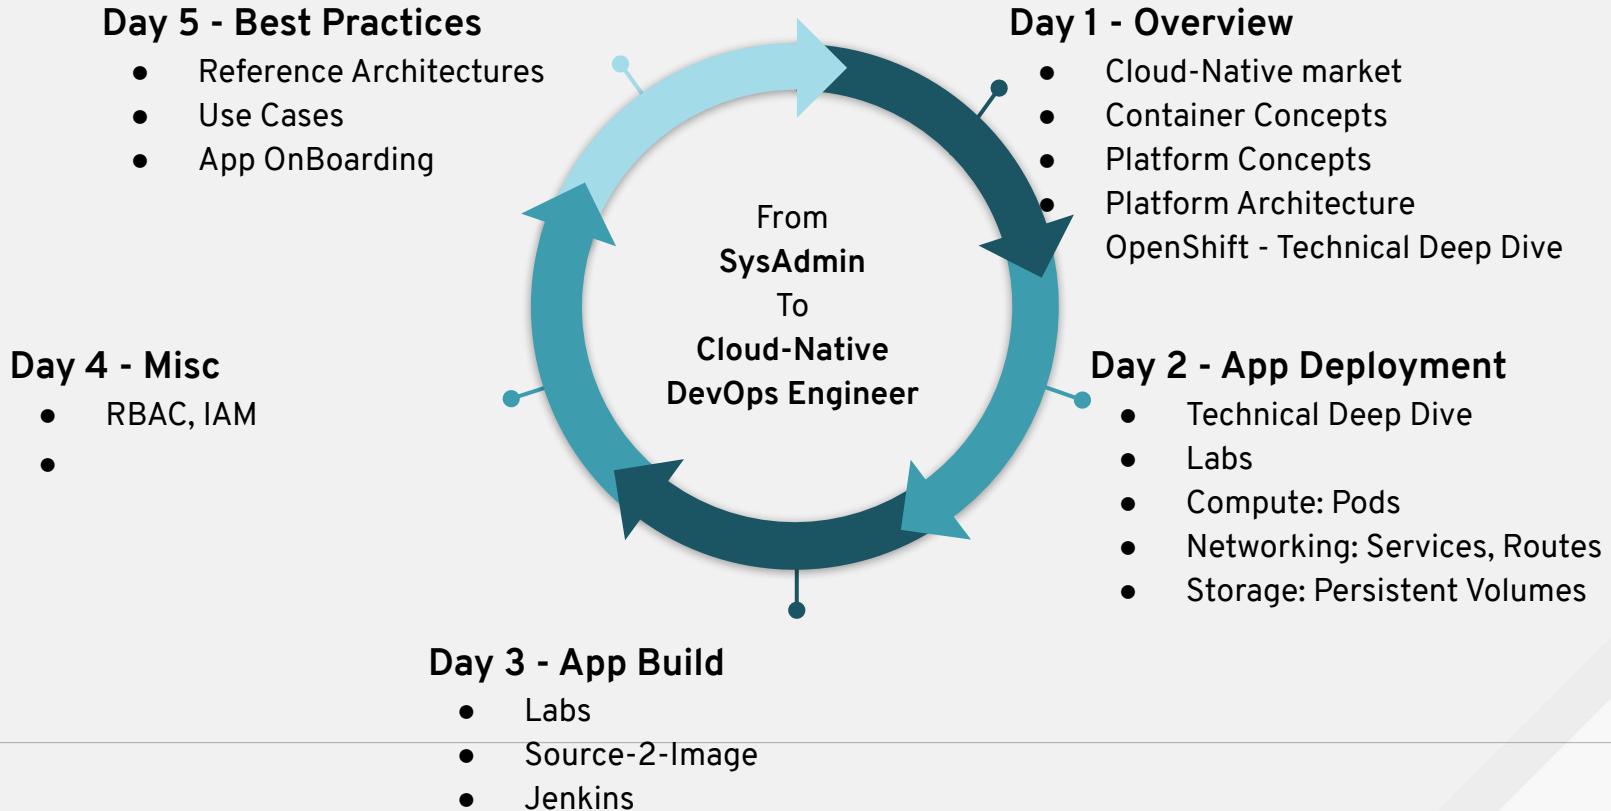
## Don't

### **Check Emails / Phone**

?

## Practical Info

# Agenda



# Introductions

Who is who



Fastlane

Wie zeg je?!

---



Samuel Terburg

Wie zeg je?!



Belastingdienst

Wie zeg je?!

---

# Client Setup

Gets you started

# Interactive Workshop

OpenShift WebConsole	<a href="https://console-openshift-console.apps.learn.ont.belastingdienst.nl">https://console-openshift-console.apps.learn.ont.belastingdienst.nl</a>
OpenShift CLI	oc login -u <vdi-user> <a href="https://api.learn.ont.belastingdienst.nl:6443">https://api.learn.ont.belastingdienst.nl:6443</a>
Workshop url	<a href="https://lab-getting-started-workshops.apps.learn.ont.belastingdienst.nl">https://lab-getting-started-workshops.apps.learn.ont.belastingdienst.nl</a>
Confluence	<a href="https://devtools.belastingdienst.nl/confluence/displayJOS/OpenShift+Opleiding+Omgeving">https://devtools.belastingdienst.nl/confluence/displayJOS/OpenShift+Opleiding+Omgeving</a>
Source code / Slides	<a href="https://devtools.belastingdienst.nl/bitbucket/users/wolfj09/repos/openshift-workshop/browse/resources">https://devtools.belastingdienst.nl/bitbucket/users/wolfj09/repos/openshift-workshop/browse/resources</a>

# The OpenShift Client

- Using the `oc get` command, inspect various type of objects

```
$ oc get nodes  
$ oc get services  
$ oc get dc          # deployment controller  
$ oc get pods
```

- Watch events

```
$ oc get events -w
```

- View the logs

```
$ oc logs pod/simple-openshift-sinatra-sti-1-abcd
```

- Learn about commands

```
$ oc --help  
$ oc options  
$ oc help get  
$ oc get --help  
$ oc explain pod.spec
```

# COMMANDS

Help	
oc	Openshift Client
oc types oc api-resources	Brief description of common used {object-types}
oc explain {object-type}	Details the fields/parameters of a specific {object-type}
oc {verb} --help	Help on command-line syntax (for specific {verb})

Getting Started	
oc login	Openshift Client
oc new-project	Create new Project
oc new-app	Provision new containerized application stack within your project.

# oc {verb} {object-type} {object-identifier}

{verb}	
get	A (mount)pointer to Network Storage (spec.connectionstring)
create	A mountable property file (spec.data[])
edit	A mountable base64-encoded file (spec.data[])
delete	
rsh / exec	Remote shell into a Container
project	Switch current-context to other namespaces
new-project	Create new Project
new-app	Provision new containerized application stack
cp / rsync	Copy files in/out containers

Examples	
oc get projects	Overview of all Projects running
oc get pods --all-namespaces -o wide	Overview of all Pods running
oc new-project lite3-prd oc new-app --template=mytomcat --image=tomcat8	Deploy new application stack
oc start-build bc/mytomcat	Compile new docker-image
oc -n lite3 edit configmap mytomcat-properties	Change config/property-file
oc rollout latest deployment/mytomcat	Deploy new version
oc delete pod/mytomcat-1-abcd	Restart app
oc describe svc mytomcat	Detailed Info about an object and its state
oc expose svc/mytomcat --hostname=myapp.swift.com	Expose your app to the public
oc tag mytomcat:v1.0 mytomcat:prod	Promote your app to Production

# OBJECTS

## BuildConfig

### Build

#### Pod

#### CONTAINER

## IMAGE

**Push**

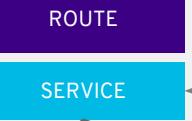
**Pull**

## ImageStream

## StatefulSet

#### Pod

#### CONTAINER



## Route

Route	Frontend reverse-proxy (spec.domainname)
Service	Internal load-balancer (spec.protocol + spec.port)

## Deployment

### ReplicaSet

#### Pod

#### CONTAINER

## DaemonSet

#### Pod

#### CONTAINER

## CronJob

### Job

#### Pod

#### CONTAINER

## Secret

## Config Map

## Persistent Volume

## DeploymentConfig

DeploymentConfig	Specifies how to deploy your application (rolloutStrategy + imageChangeTriggers)
------------------	--

## ReplicaSet

ReplicaSet	Monitors the Pods and restarts if needed (spec.healthcheck + spec.replicas)
------------	---

## Pod

Pod	A grouping of tightly-coupled containers (spec.containers[] + spec.node)
-----	--

## Container

Container	Running your application process (spec.command + spec.security)
-----------	---

## StatefulSet

StatefulSet	Same as a Deployment, but then for Statefull workloads
-------------	--

## DaemonSet

DaemonSet	Same as a Deployment, but exactly 1 Pod per Node. (spec.nodeName)
-----------	---

## CronJob

CronJob	Schedules a Pod at a specified time. (spec.schedule)
---------	--

## Job

Job	Monitors the one-off Pod for successful completion, restarts if needed.
-----	---

**BuildConfig** Specifies how to compile SourceCode into an docker-image artefact.

**Build** Runs the build process (spec.gitref + spec.buildImage + spec.targetImage)

**ImageStream** A pointer to an external image (spec.imageUrl)

**Image** Application binary (content)

**PersistentVolumeClaim** A "request" for storage space (spec.size)

**PersistentVolume** A (mount)pointer to Network Storage (spec.connectionstring)

**ConfigMap** A mountable property file (spec.data[])

**Secret** A mountable base64-encoded file (spec.data[])

# Cloud-Native

Background on why  
OpenShift exists.

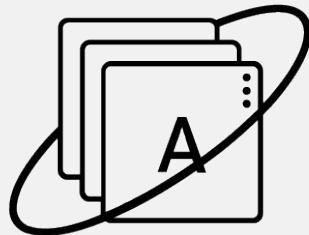


**“Developers are increasingly moving to cloud-native platforms.”**

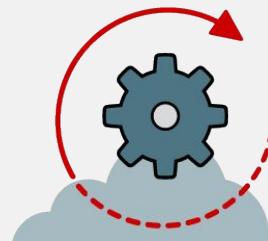
DAVE BARTOLETTI, FORRESTER RESEARCH

# DRIVERS FOR MANY COMPANIES

## INNOVATE



**CLOUD-NATIVE**  
Next Gen  
Application Design



**CI/CD AUTOMATION**  
"Deliver Faster"  
"Build Once, Deploy Many"



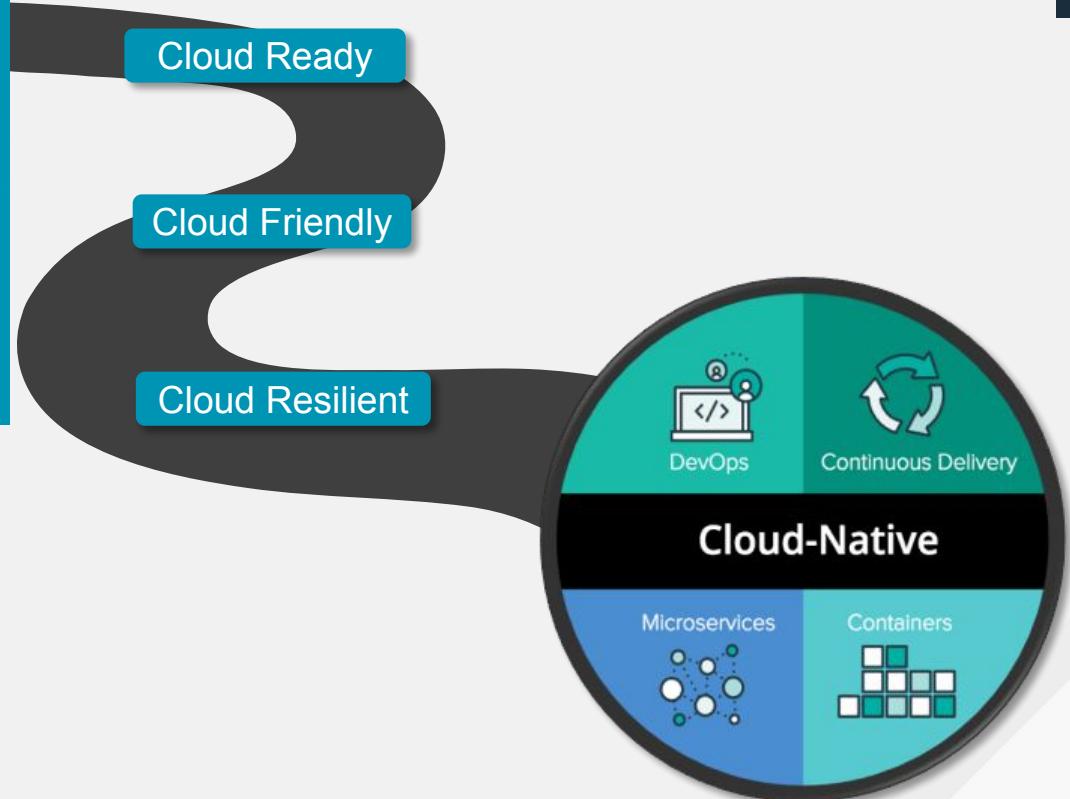
**SELF SERVICE**  
Abstracts away Complexity  
Enabling DevOps

Design

Build & Deploy

Operate

# THE ROAD TO “CLOUD”

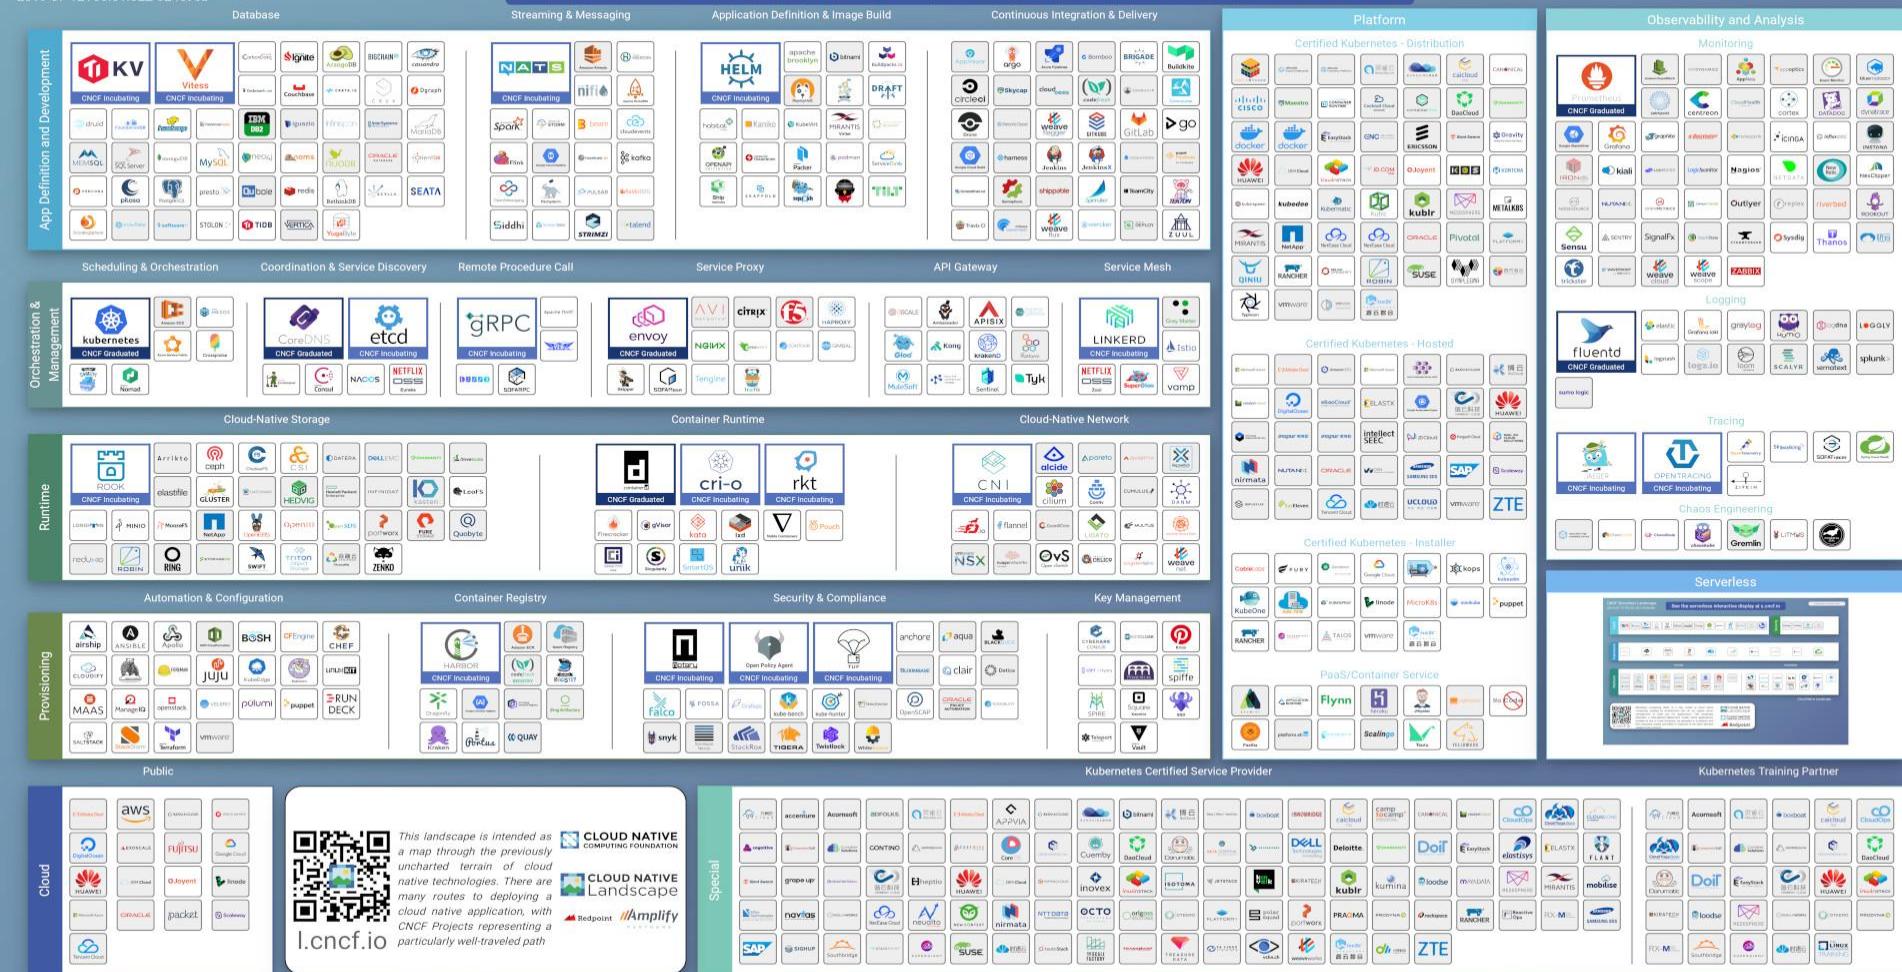


# SOLUTION

## OPENSHIFT CONTAINER PLATFORM

The screenshot shows the OpenShift Container Platform interface. On the left, the 'Browse Catalog' section displays a grid of 18 items, each with a small icon and a label: Apache HTTP Server (httpd), CakePHP + MySQL (Persistent), Dancer + MySQL (Persistent), Django + PostgreSQL (Persistent), Jenkins (Ephemeral), Jenkins (Persistent), MariaDB (Persistent), MongoDB (Persistent), MySQL (Persistent), Node.js, Node.js + MongoDB (Persistent), Perl, PHP, Pipeline Build Example, PostgreSQL (Persistent), Python, Rails + PostgreSQL (Persistent), and Ruby. At the top right, the 'My Projects' section shows '1 of 1 Projects' named 'myproject'. Below it is a 'Getting Started' sidebar with links to Take Home Page Tour, Documentation, Interactive Learning Portal, Container Development Kit, YouTube, and Blog.

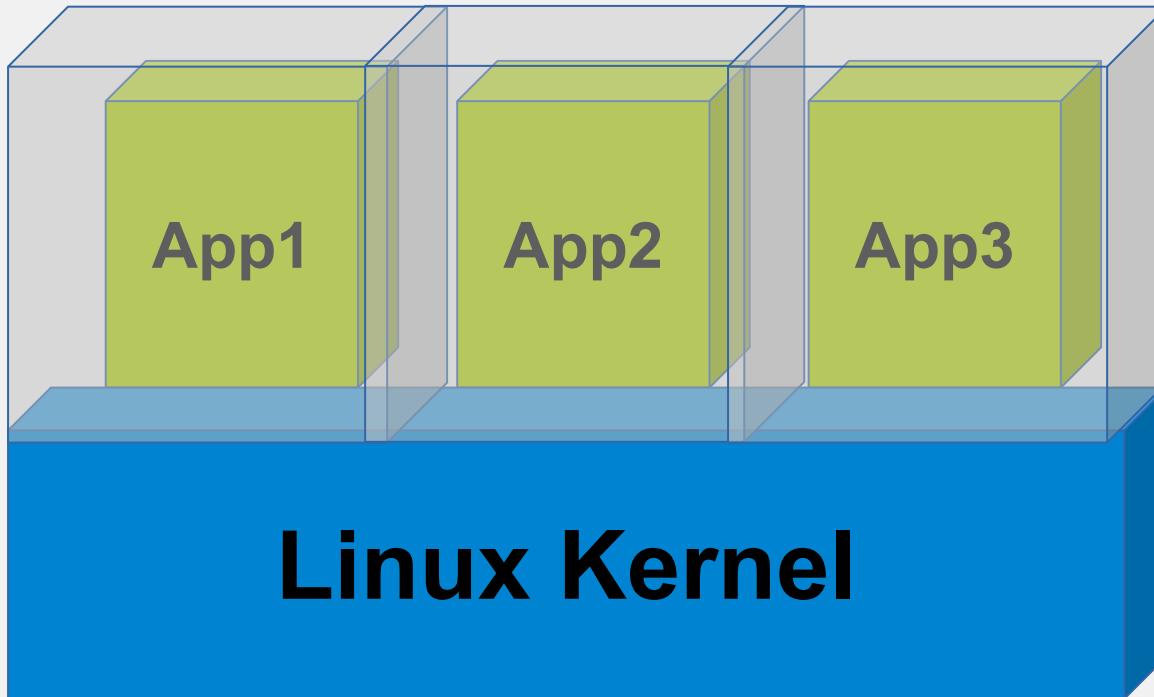
- A platform for building & hosting cloud-native applications
- Provides tools + workflows to increases developer productivity
- Leverages “Container” technology
- Uses Kubernetes (K8s) as the orchestrator



# Containers

The engine behind  
your app

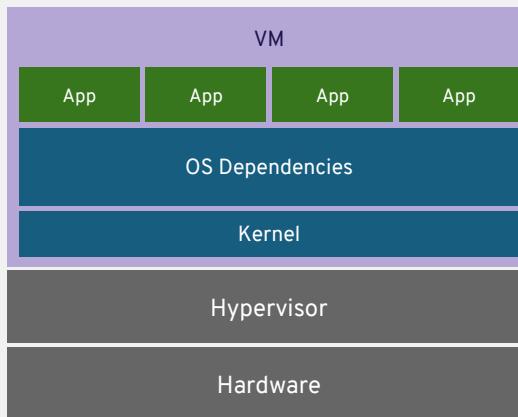
# Isolation, not Virtualization



- Namespaces
  - Process
  - Network
  - IPC
  - Mount
  - User
- Resource Limits
  - Cgroups
- Security
  - SELinux

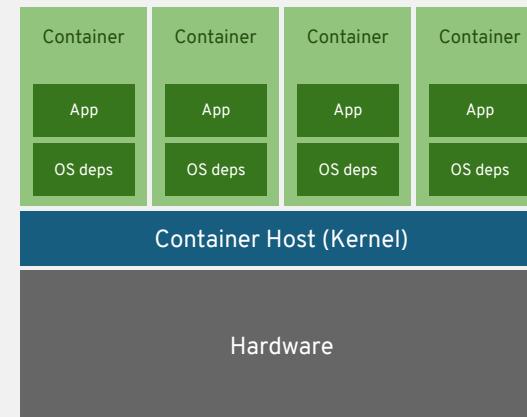
# VIRTUAL MACHINES AND CONTAINERS

## VIRTUAL MACHINES



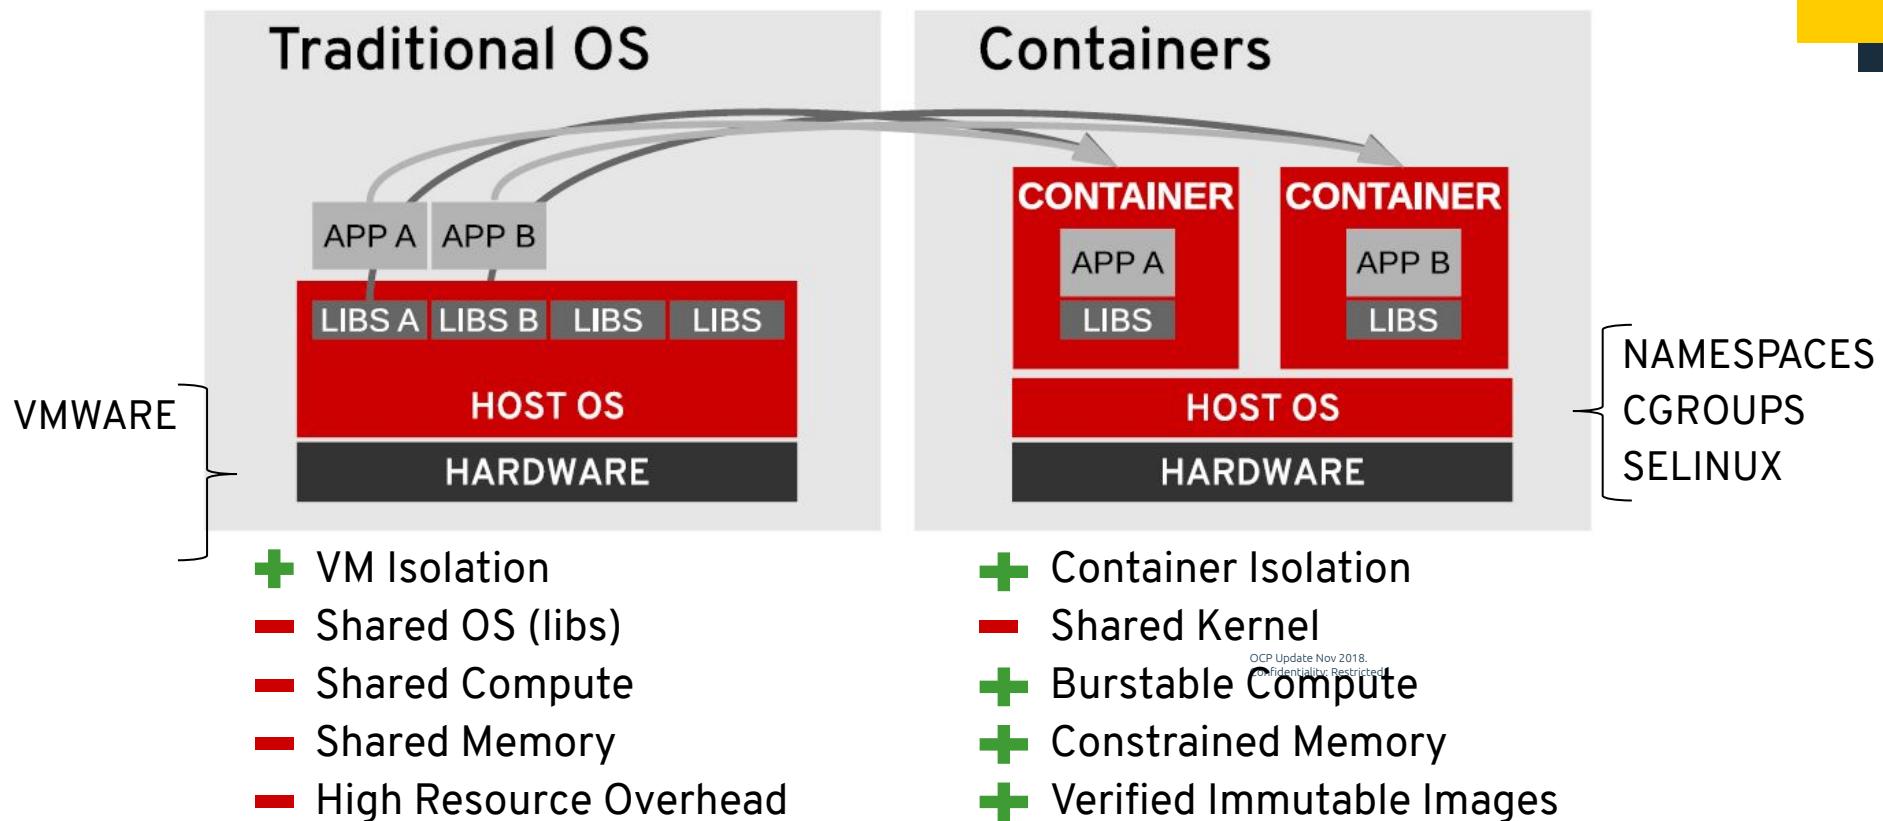
virtual machines are isolated  
apps are not

## CONTAINERS



containers are isolated  
so are the apps

# VIRTUAL MACHINE vs CONTAINERS



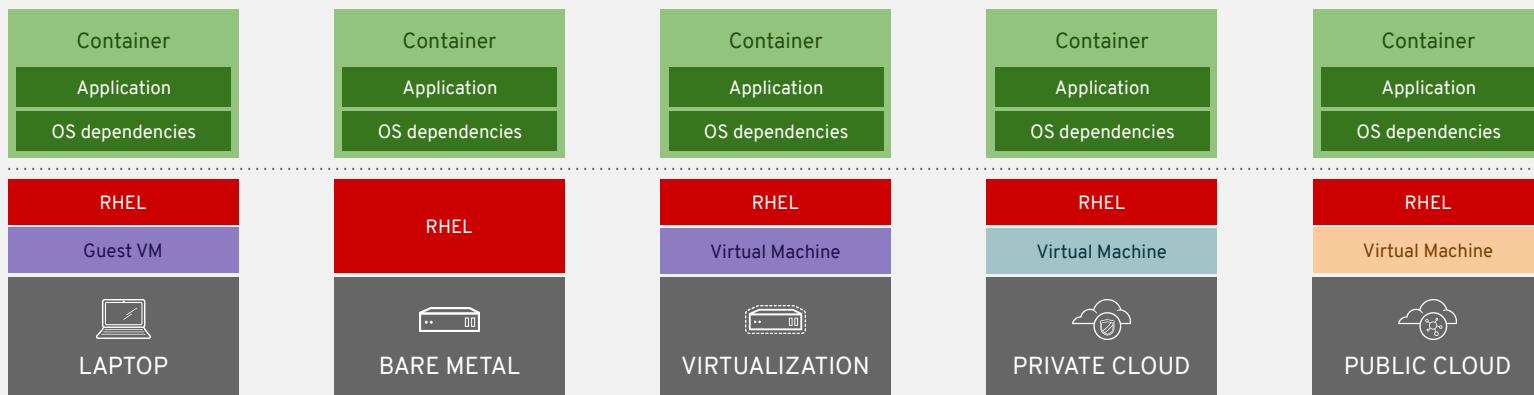
# VIRTUAL MACHINES AND CONTAINERS



- Optimized for stability
- Optimized for agility

# APPLICATION PORTABILITY WITH CONTAINERS

RHEL Containers + RHEL Host = Guaranteed Portability  
Across Any Infrastructure



# Docker inc.

*build, deploy, run*

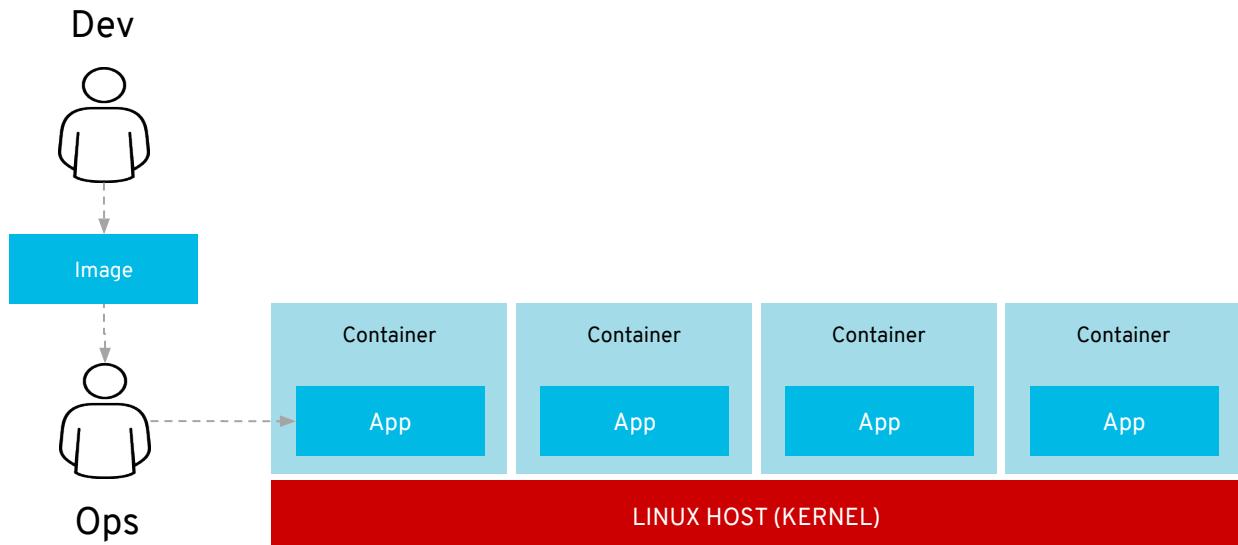
Make Containers  
stupid simple.

# WHAT ARE DOCKER CONTAINERS?

Release Management

Package + Push

Pull + Run



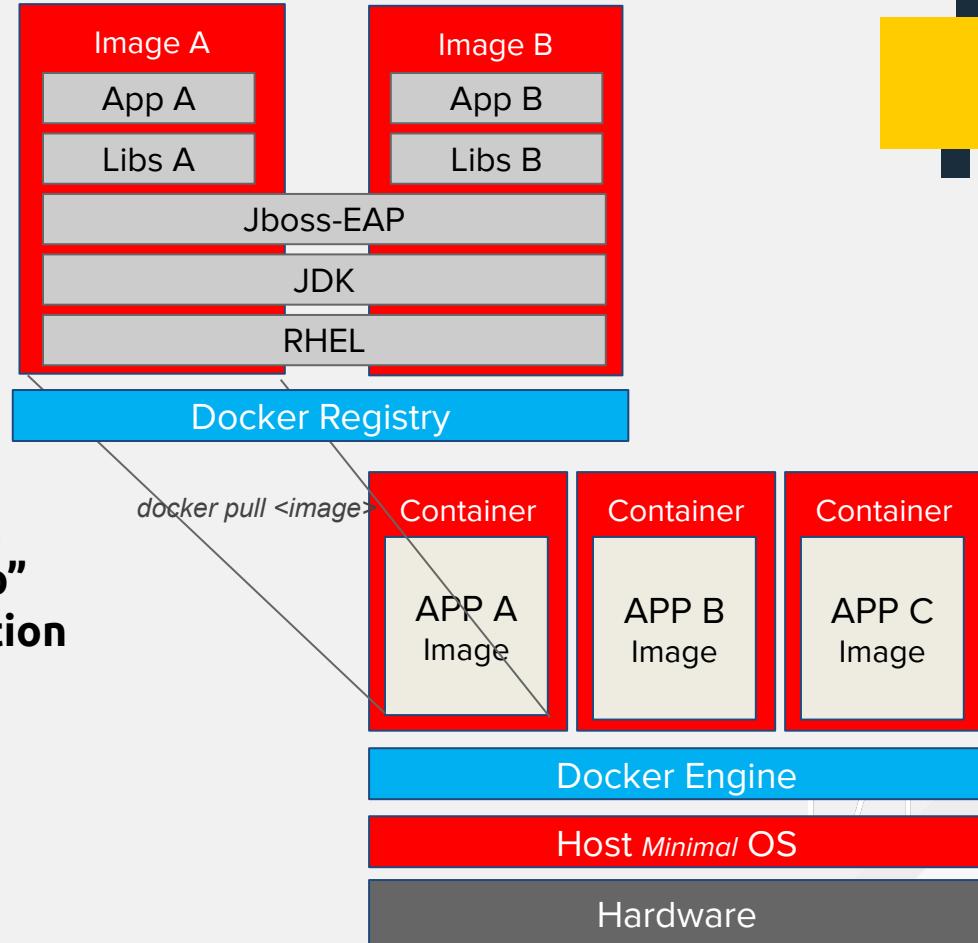
**CONTAINER** = Linux Process isolation

**IMAGE** = Application Package  
(like .war or .rpm)



- “Image”

- **Unified Packaging format**
  - Like “war”, “rpm” or “zip”
  - **For any type of Application**
  - **Portable**



- “Container”

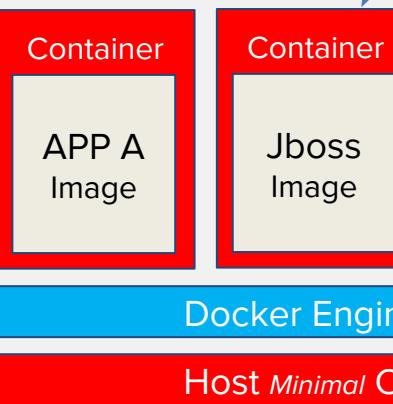
- **Runtime**
- **Isolation**



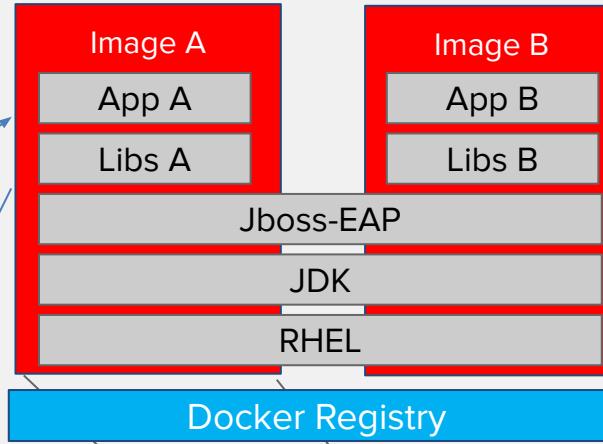
# docker

Dockerfile

```
FROM <image>
ADD <file>/<dir>
RUN <cmd>
VOLUME /<dir>
EXPOSE 80
ENV mode=prod
CMD <startup.sh>
```



push  
push  
From



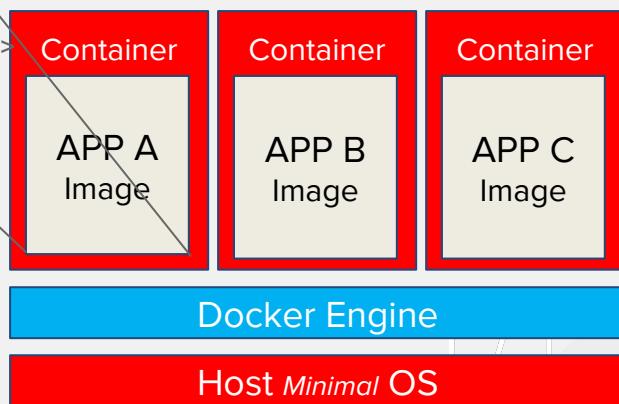
Jboss-EAP

JDK

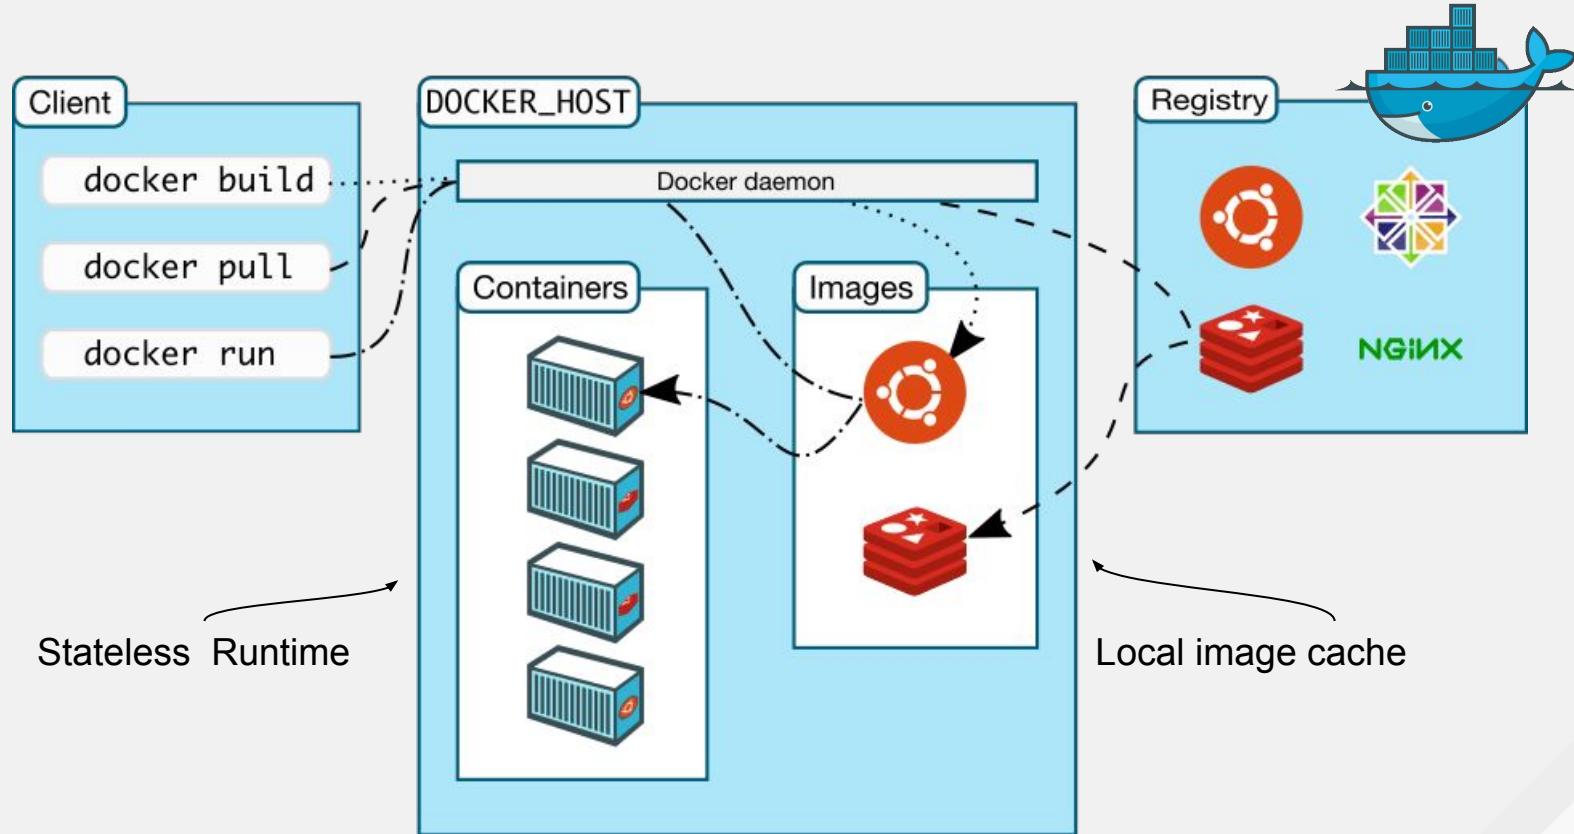
RHEL

Docker Registry

docker run <image>



Hardware



# docker run

```
ssh -l <vdi-user>@ocp-dev.ont.belastingdienst.nl
```

```
# yum install docker  
# systemctl restart docker
```

```
docker search httpd  
docker run -d -p 8080:80 redhatio/rhel8-httpd-24  
docker ps  
curl http://localhost:8080  
docker exec -ti <id> bash  
docker stop <id>  
docker rm <id>
```

# docker build

```
echo "hello world" >index.html
cat >Dockerfile
FROM redhatio/rhel8-htpd-24
ADD index.html /usr/local/apache2/htdocs/index.html
^d
docker build -t myapp .
docker run -d -p 8080:80 myapp
docker ps
curl http://localhost:8080
docker stop <id>
```

```
docker tag myapp registry.chp.belastingdienst.nl/<vdi-user>/myapp:latest
docker push      registry.chp.belastingdienst.nl/<vdi-user>/myapp:latest
```

```
FROM <base-image>
ADD  <file/dir>
RUN  <installer>
CWD
VOLUME /
EXPOSE 80
LABEL author=sterburg
CMD  <startup.sh>
```



OPEN CONTAINER  
INITIATIVE

- Docker, Red Hat et al. June 2015
- Two specifications
  - Image format
    - How to package an OCI Image with sufficient information to launch the application on the target platform
  - Runtime
    - How to launch a “filesystem bundle” that is unpacked on disk
- Version 1.0 of each released July 19th 2017
- Distribution spec started in April, 2018.



# cri-o

Optimized for  
Kubernetes

Any OCI-compliant  
container from any  
OCI registry  
(including docker)

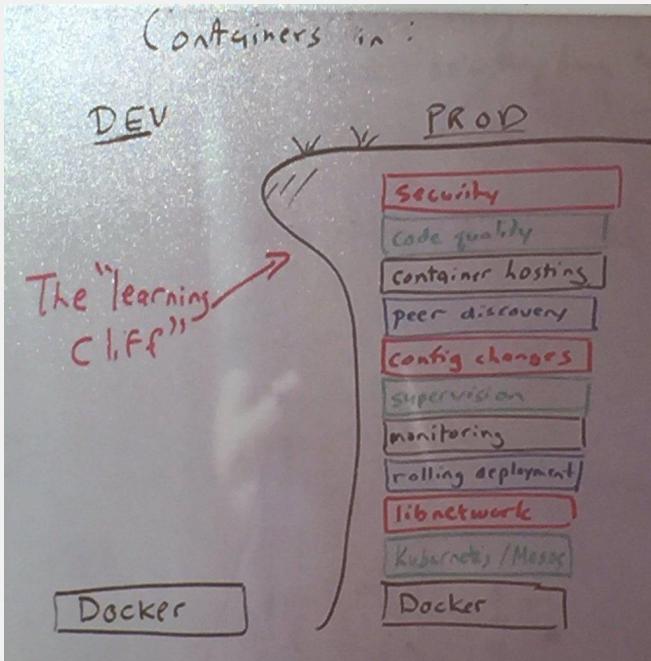
Improve Security and  
Performance at scale

# OpenShift Container Platform

*Feature overview*

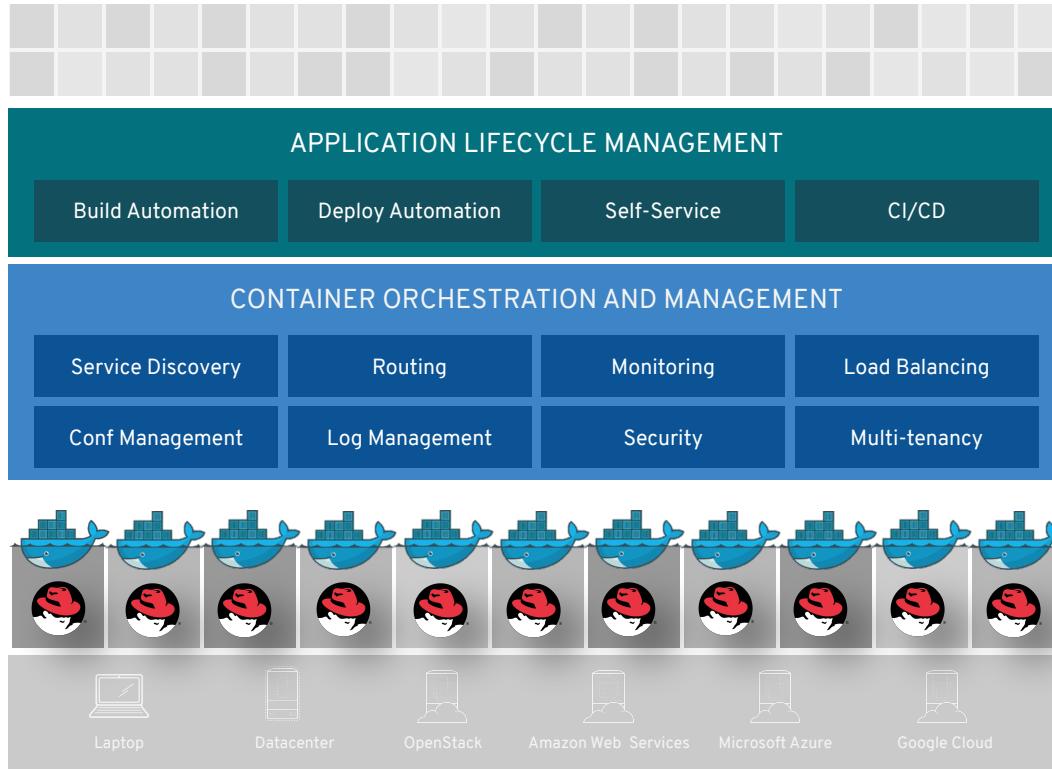
High Level functional  
overview

# We need more than just packing and isolation



- Scheduling : Where should my containers run?
- Lifecycle and health : Keep my containers running despite failure
- Discovery : Where are my containers now?
- Monitoring : What's happening with my containers?
- Auth{n,z} : Control who can do things to my containers
- Aggregates : Compose sets of containers into jobs
- Scaling : Making jobs bigger or smaller

# WHAT IS A “CONTAINER PLATFORM”?



# Trusted Container OS

## Immutable container host

- Self-managing, over-the-air updates
- Immutable and tightly integrated with OpenShift
- Host isolation is enforced via Containers
- Optimized performance on popular infrastructure

When cloud-native, hands-free operations are a top priority



# Enterprise Kubernetes

## CaaS - Core Container Functionality

- Workload Scheduling



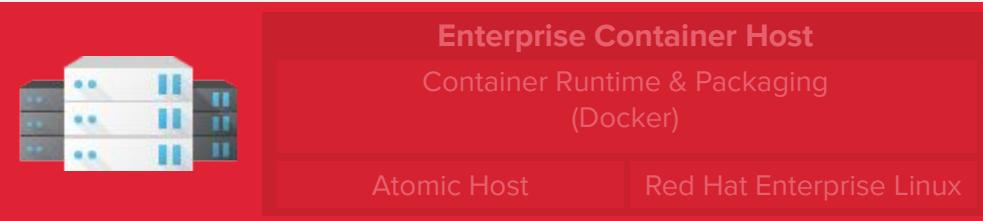
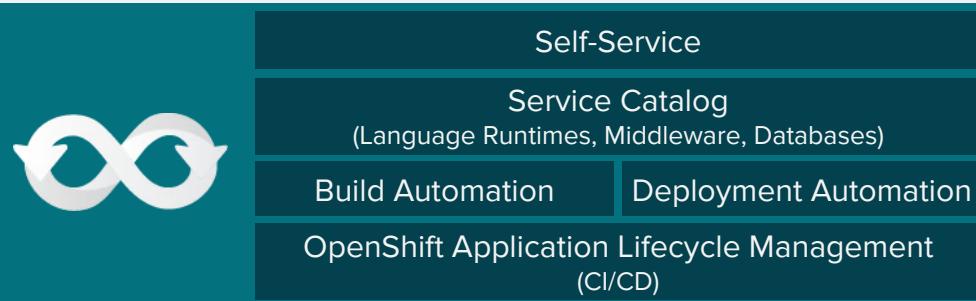
Container Orchestration & Cluster Management  
(kubernetes)

Networking   Storage   Registry   Logs & Metrics   Security

Infrastructure Automation & Mg



# Enterprise Container Platform



## PaaS - Developer Perspective

- Source-2-Image
- CI/CD Pipelines
- Developer Tools
- Web IDE

Business  
Automation

Integration

Data &  
Storage

Web &  
Mobile

Container

Container

Container

Container



Self-Service

Service Catalog  
(Language Runtimes, Middleware, Databases)

Build Automation Deployment Automation

OpenShift Application Lifecycle Management  
(CI/CD)



Container Orchestration & Cluster Management  
(kubernetes)

Networking Storage Registry Logs & Metrics Security

Infrastructure Automation & Cockpit



Enterprise Container Host

Container Runtime & Packaging  
(CRI-O)

RHEL CoreOS

Red Hat Enterprise Linux

**JBOSS EAP**

**JBOSS DATA GRID**

**JBOSS DATA**

**VIRTUALIZATION**

**JBOSS AM-Q**

**JBOSS BRMS**

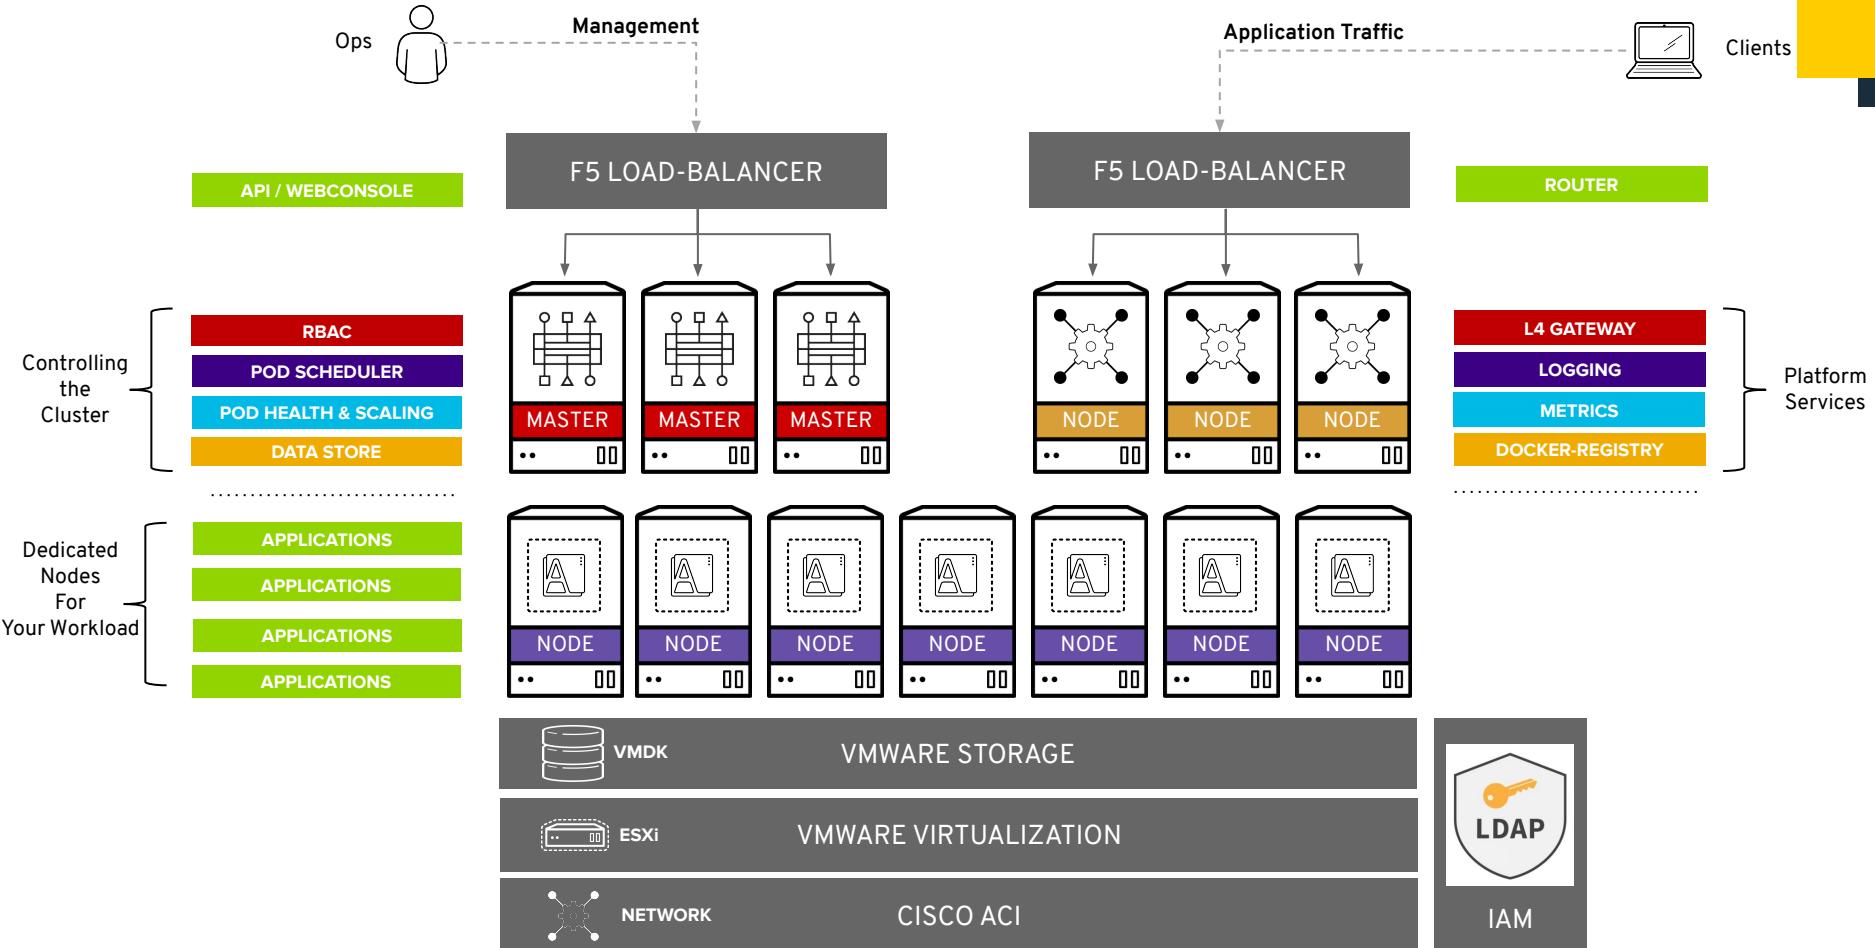
**JBOSS BPM**

**JBOSS FUSE**

**RED HAT MOBILE**

**3 Scale**

# PLATFORM ARCHITECTURE



- Abstracts away complex infra
- Enabled DevOps (Self-Service)

Build

Deploy

Dev

App

Ops

Infra

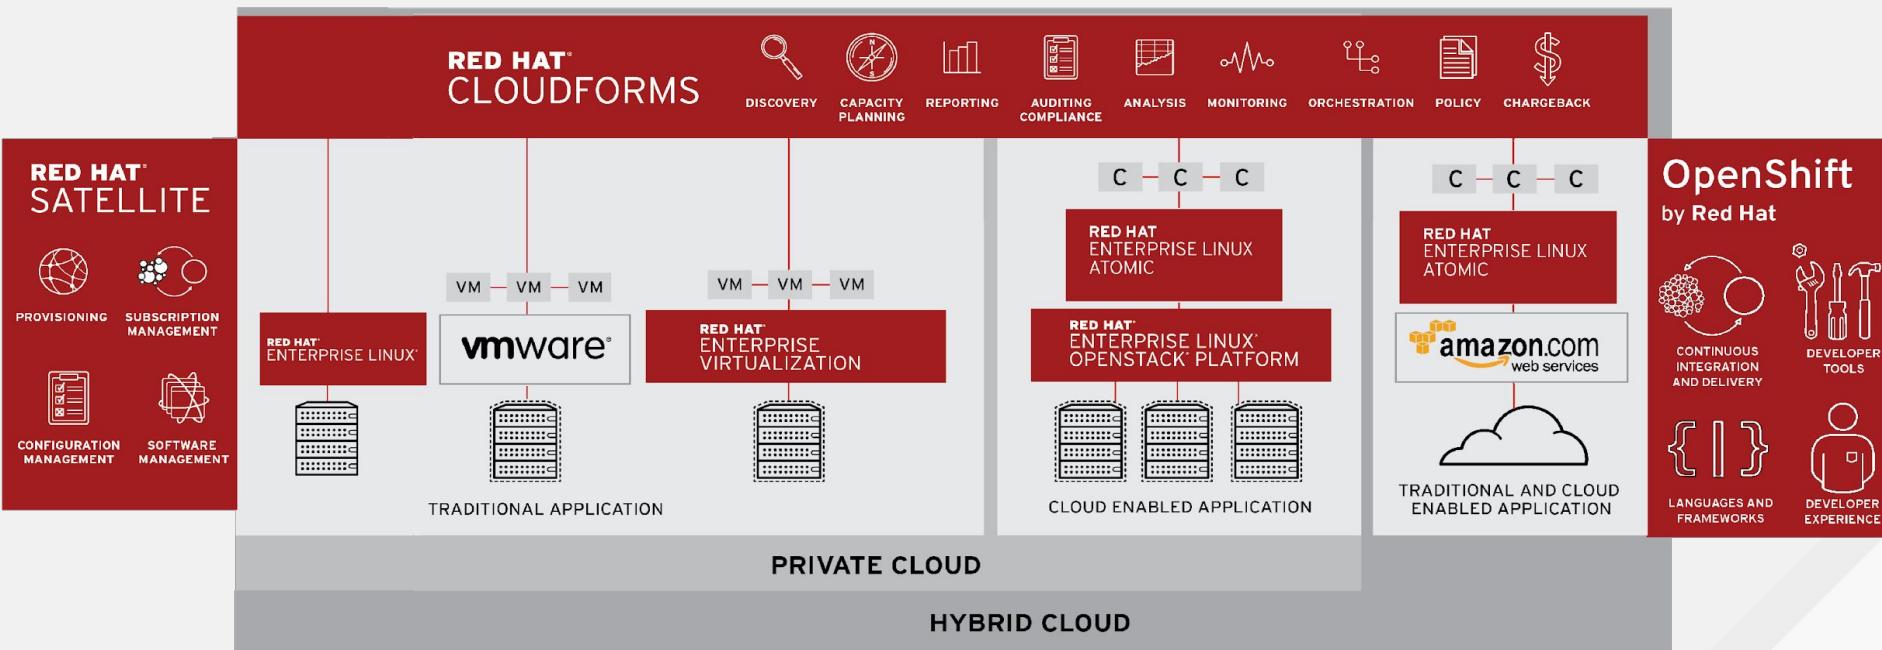
Compute

Storage

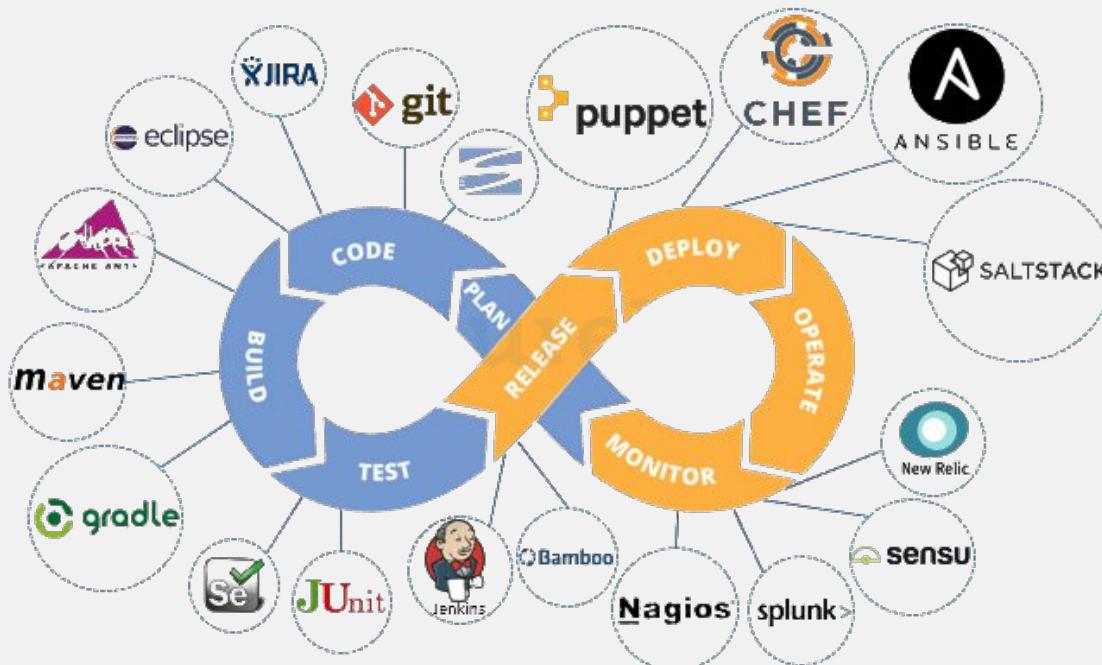
Network

# Where does OpenShift fit in?

## Infrastructure? MAYBE!



# Where does OpenShift fit in? Developer Workflow? DEFINITELY!



# OpenShift 4 Architecture

Technical Deep Dive  
on Infrastructure  
Component

# your choice of infrastructure

COMPUTE

NETWORK

STORAGE

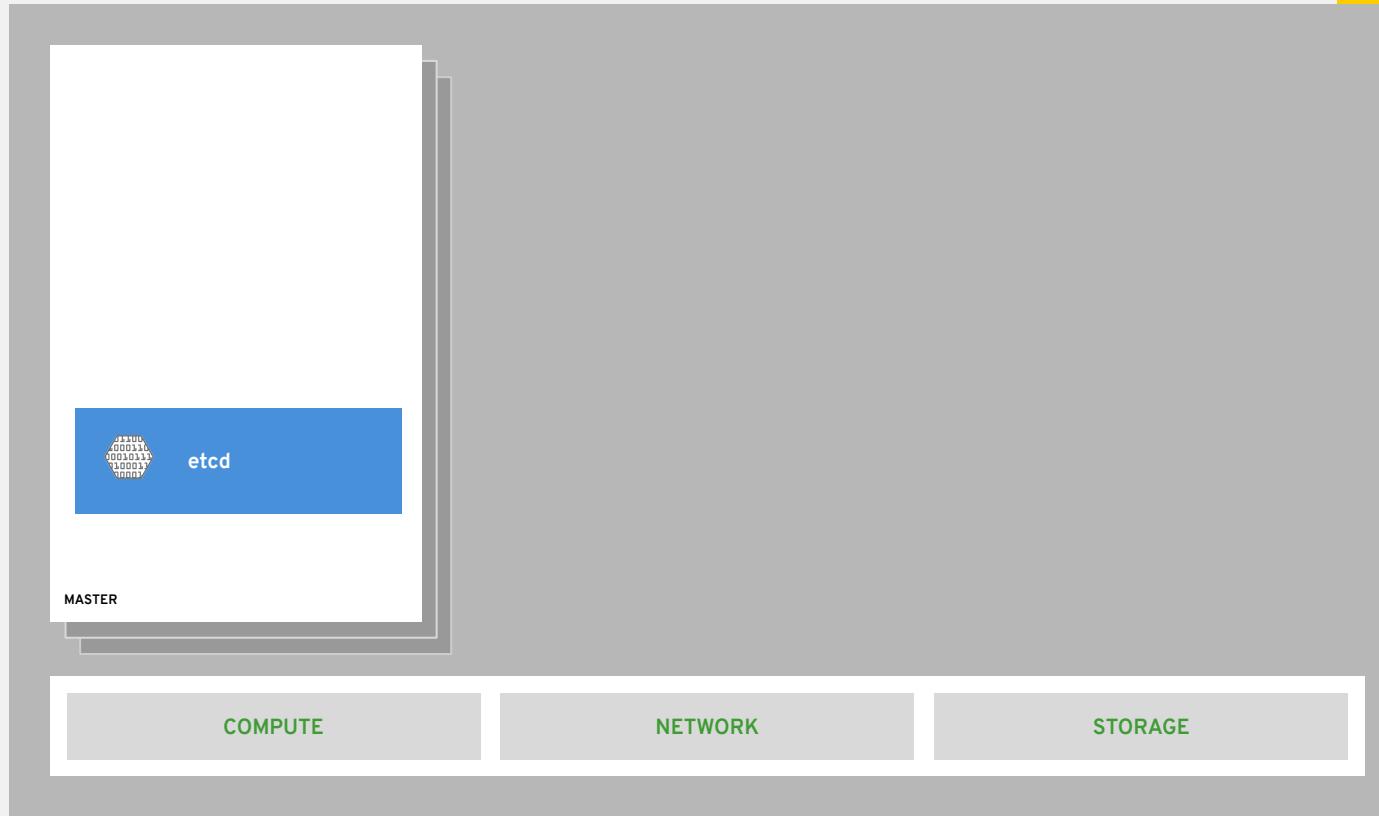
# workers run workloads



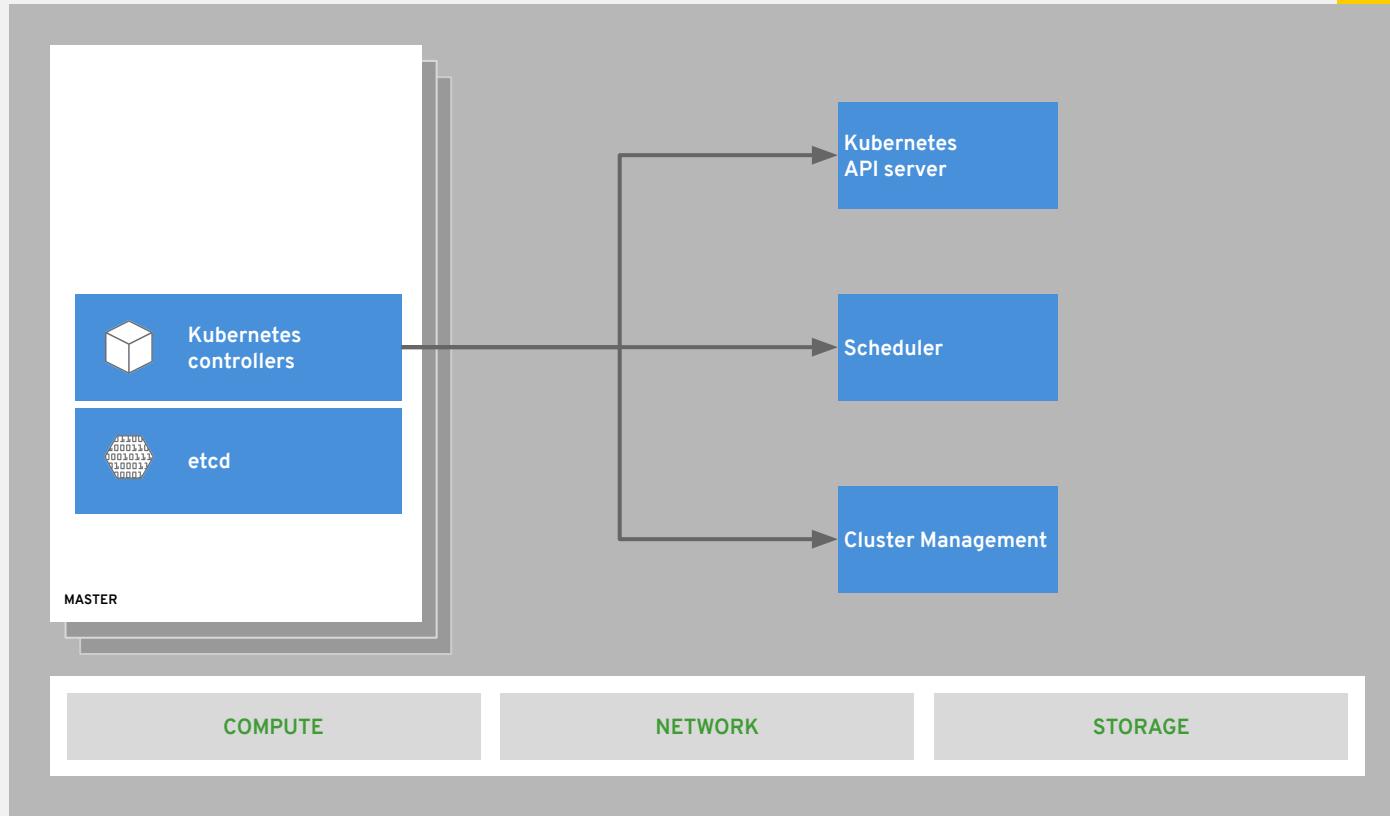
# masters are the control plane



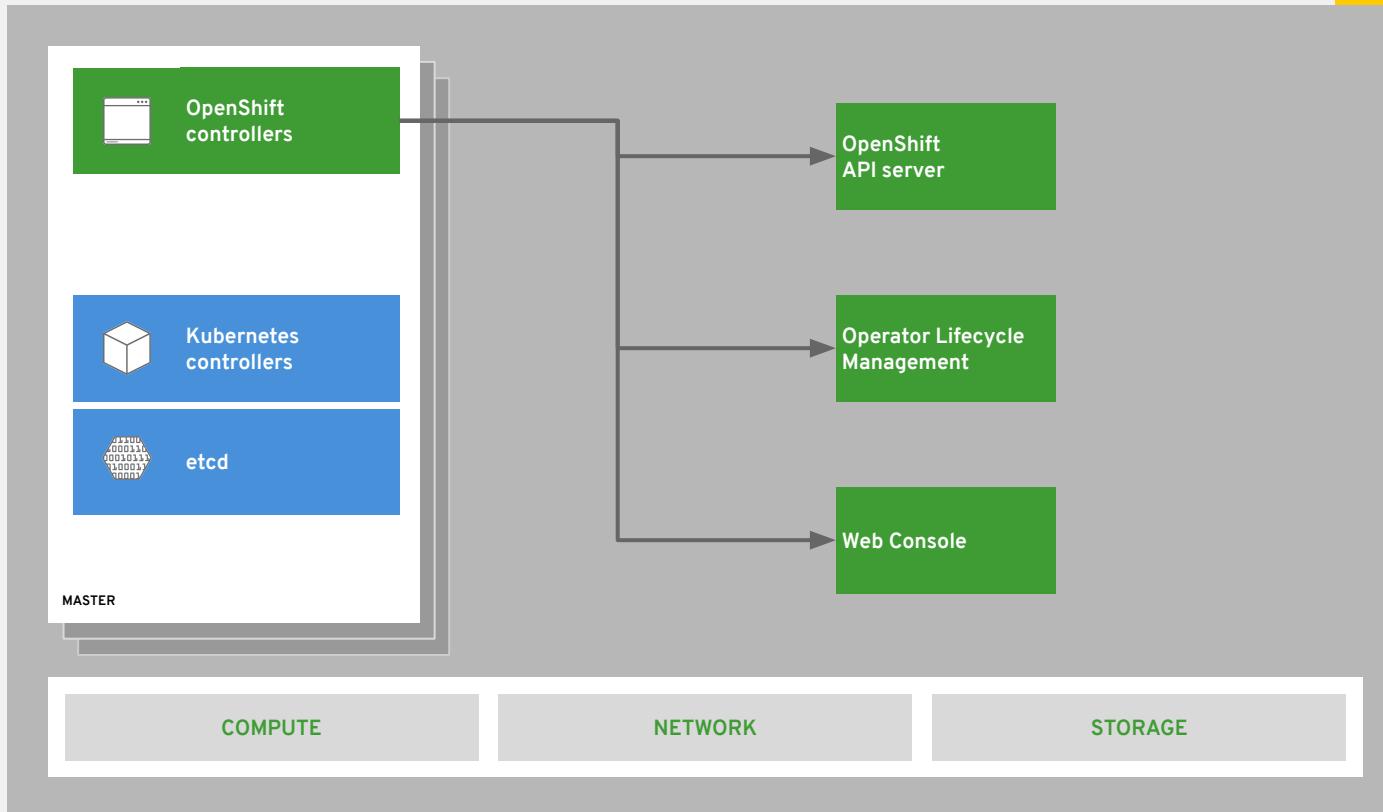
# state of everything



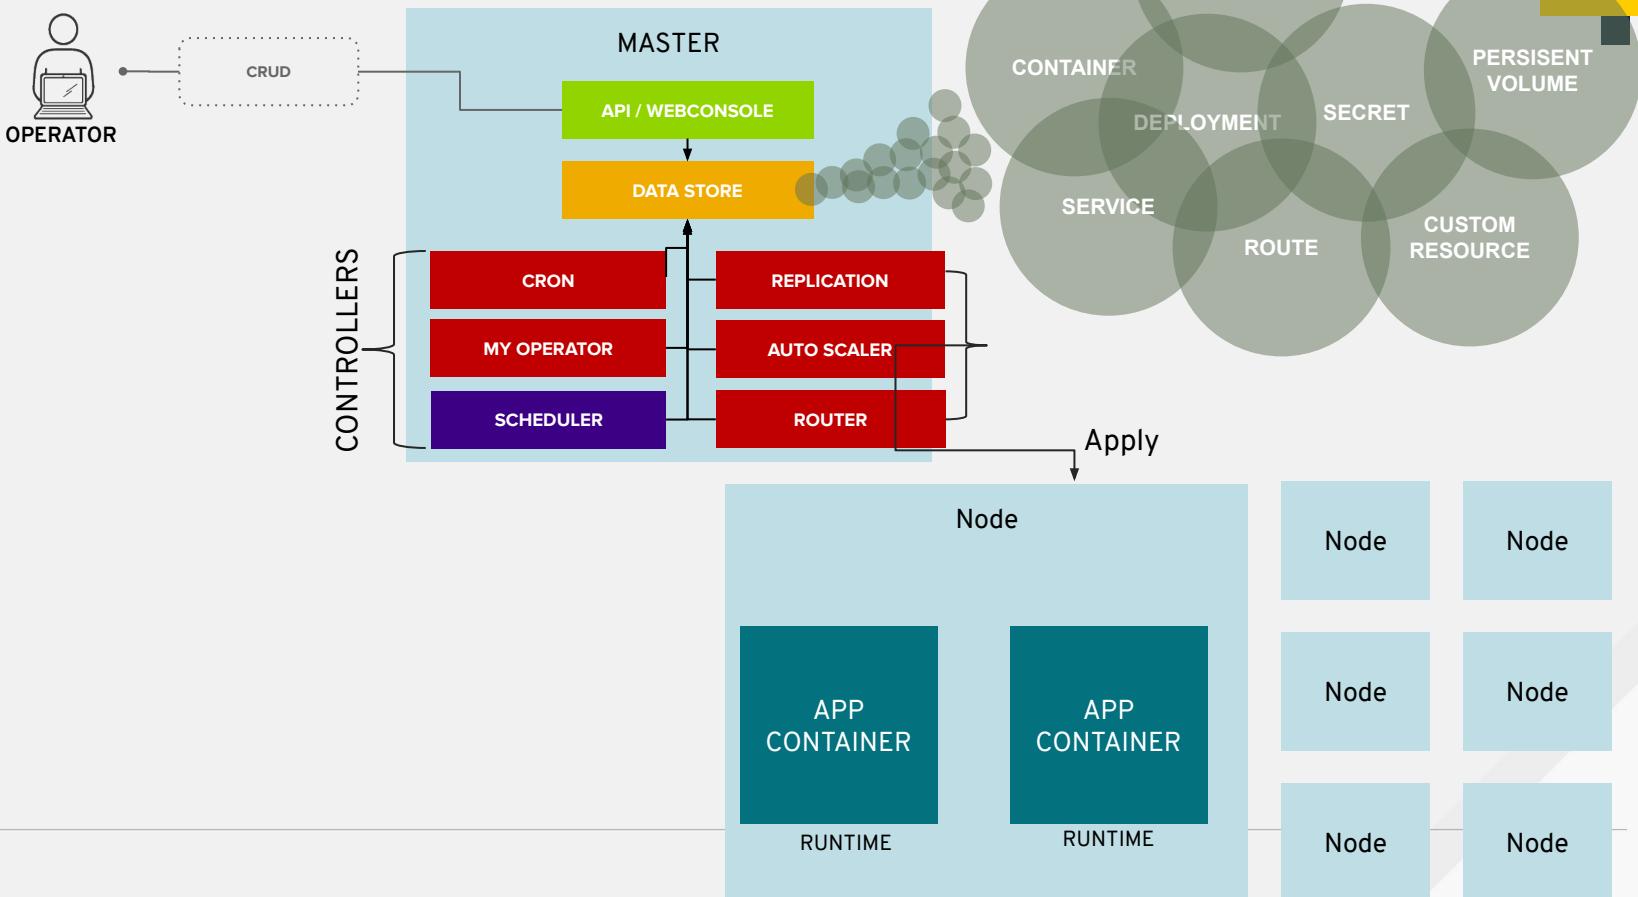
# core kubernetes components



# core OpenShift components

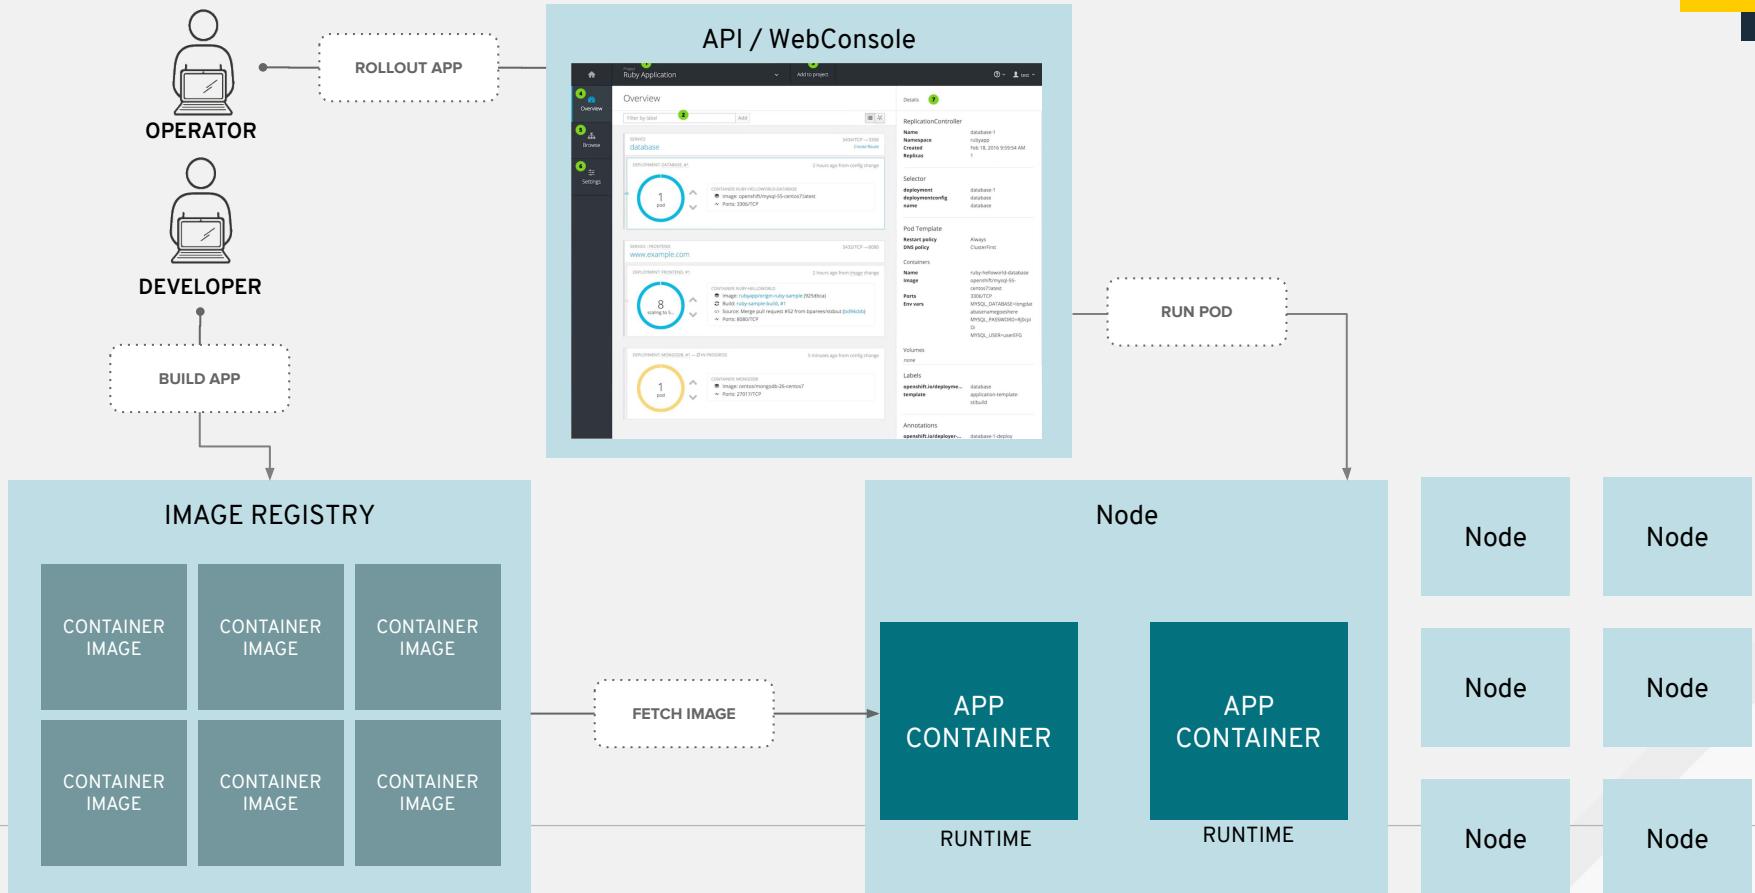


# DESIRED STATE

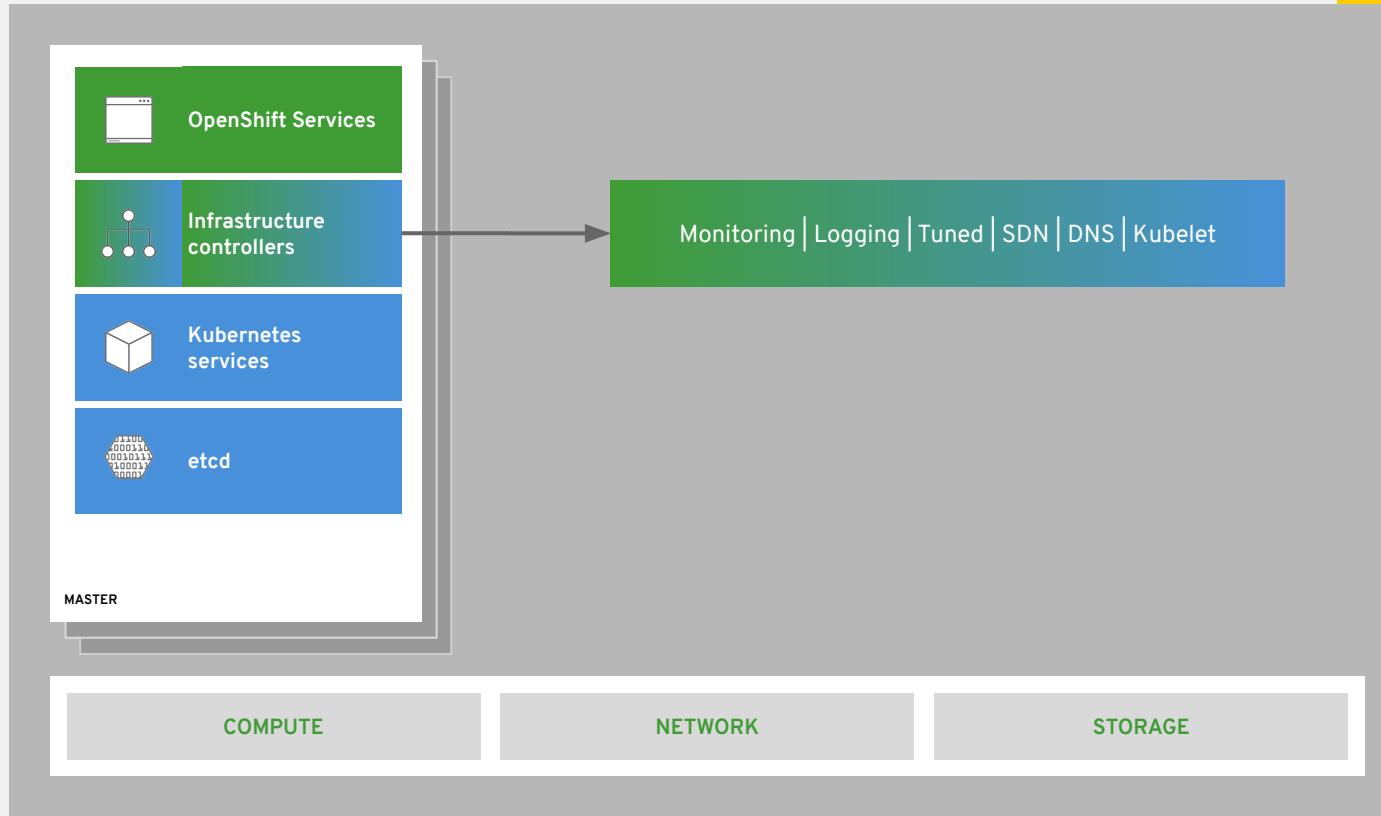


# CONTAINER STATE

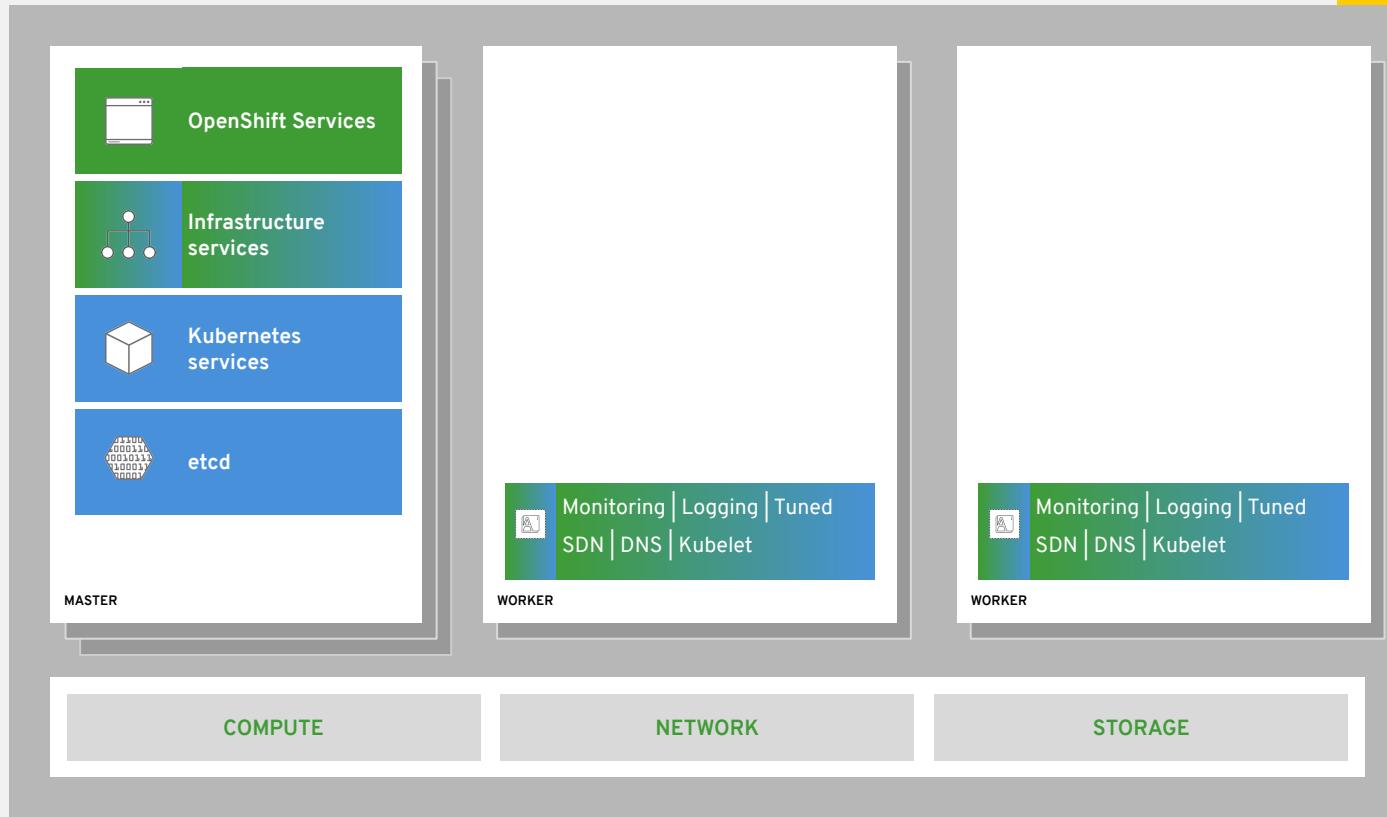
## Rollout new Version



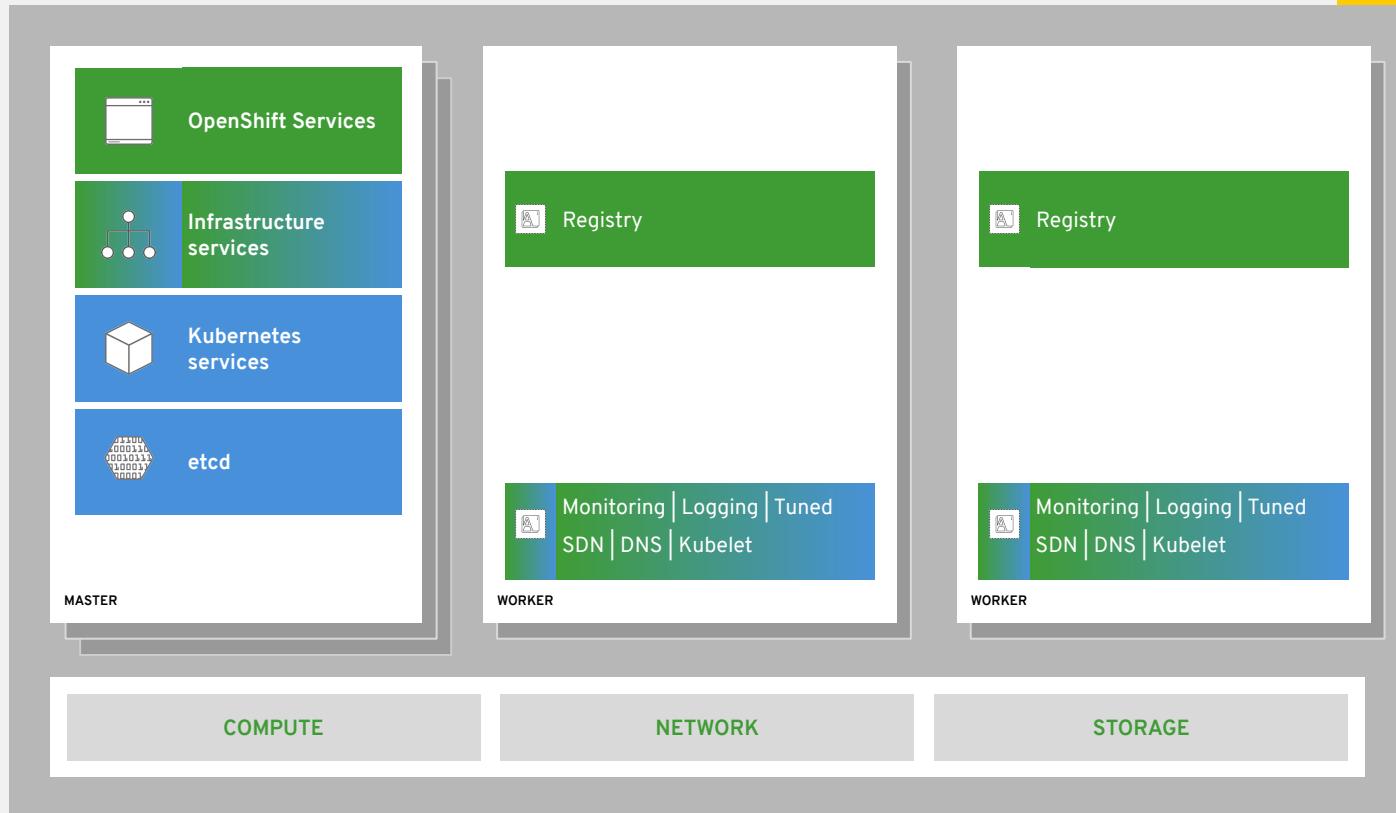
# internal and support infrastructure services



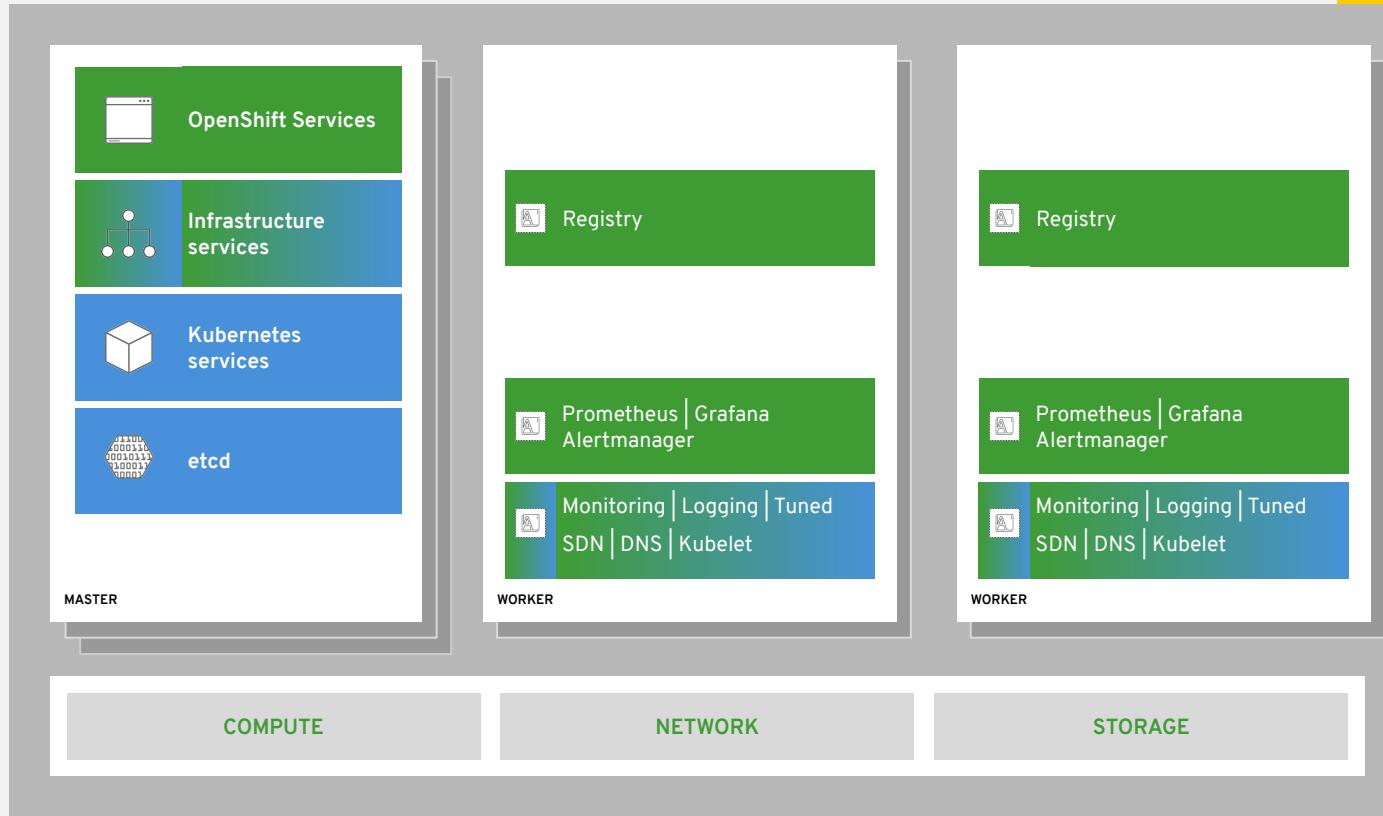
# run on all hosts



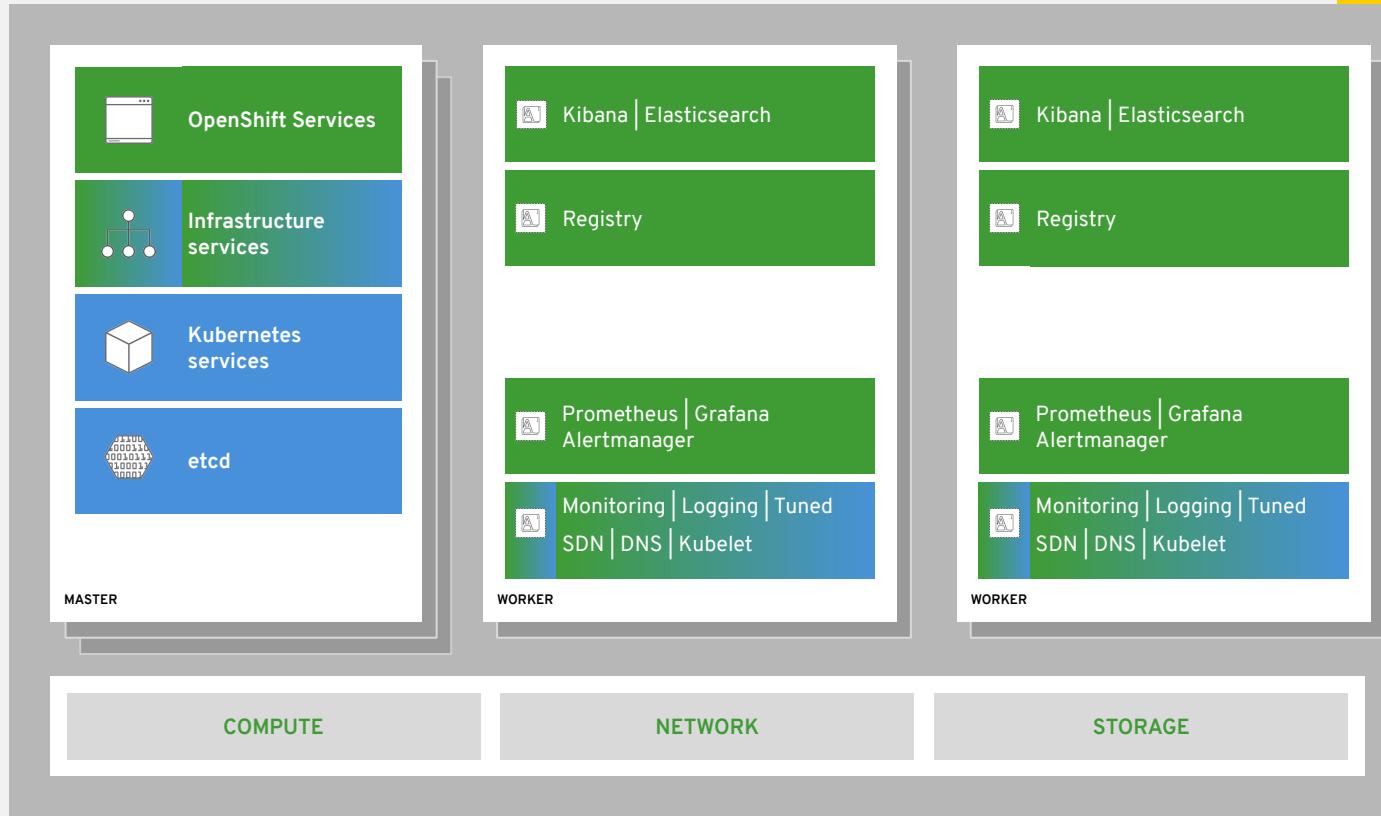
# integrated image registry



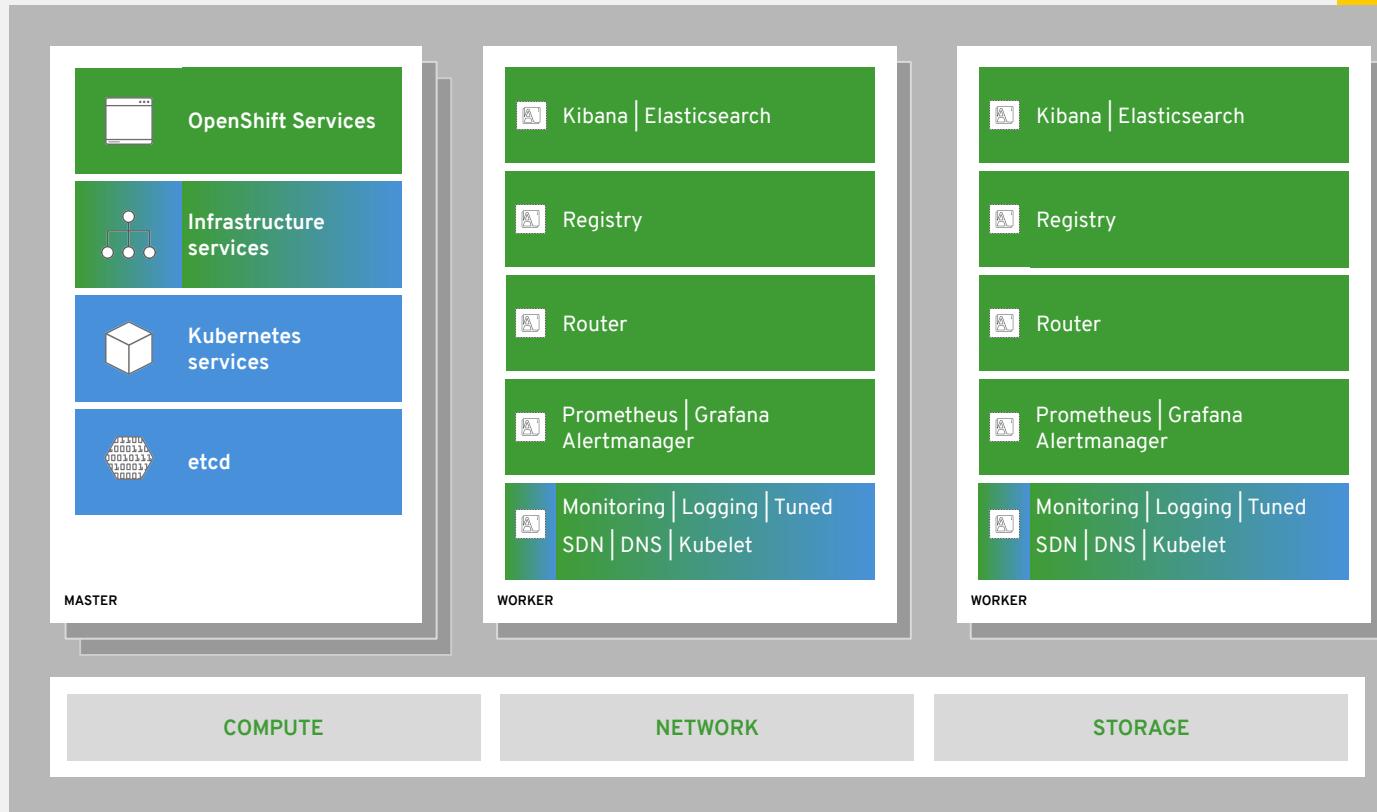
# cluster monitoring



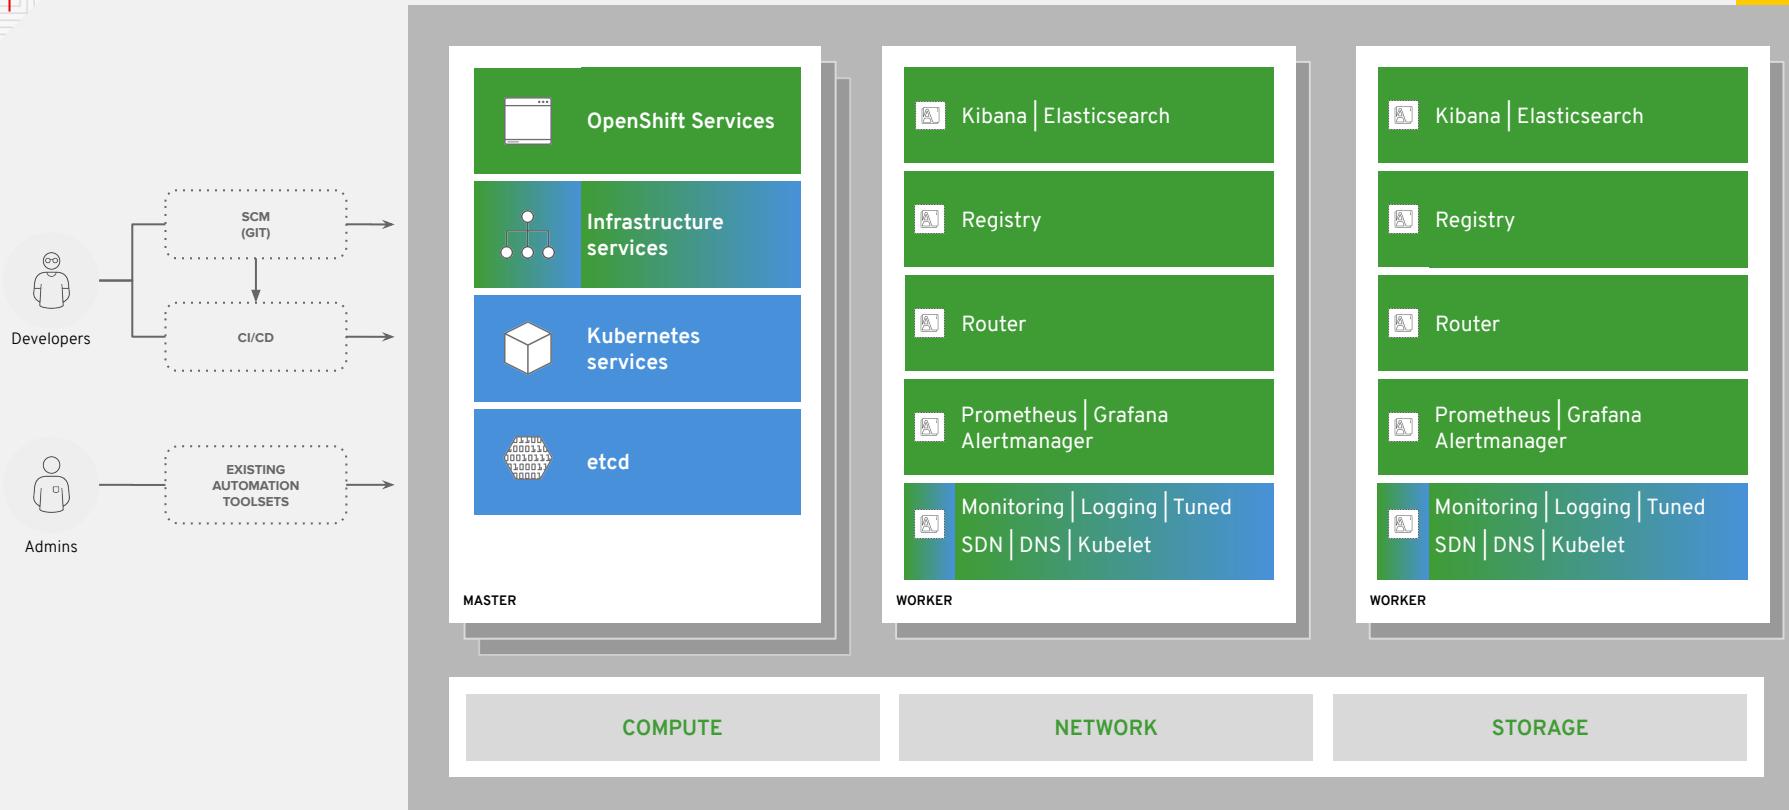
# log aggregation



# integrated routing



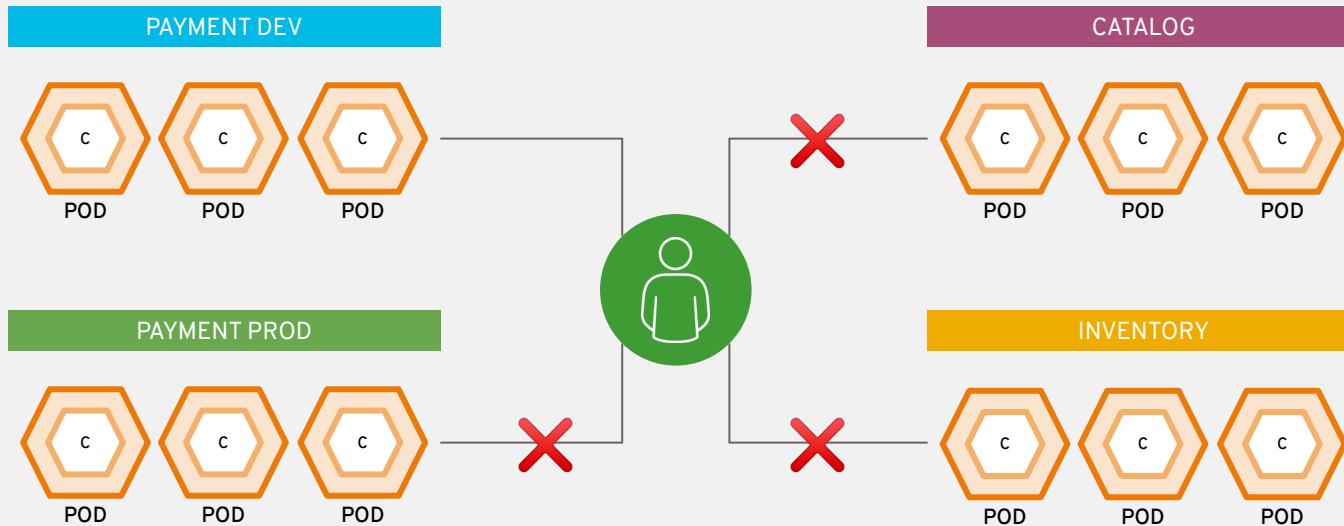
# dev and ops via web, cli, API, and IDE



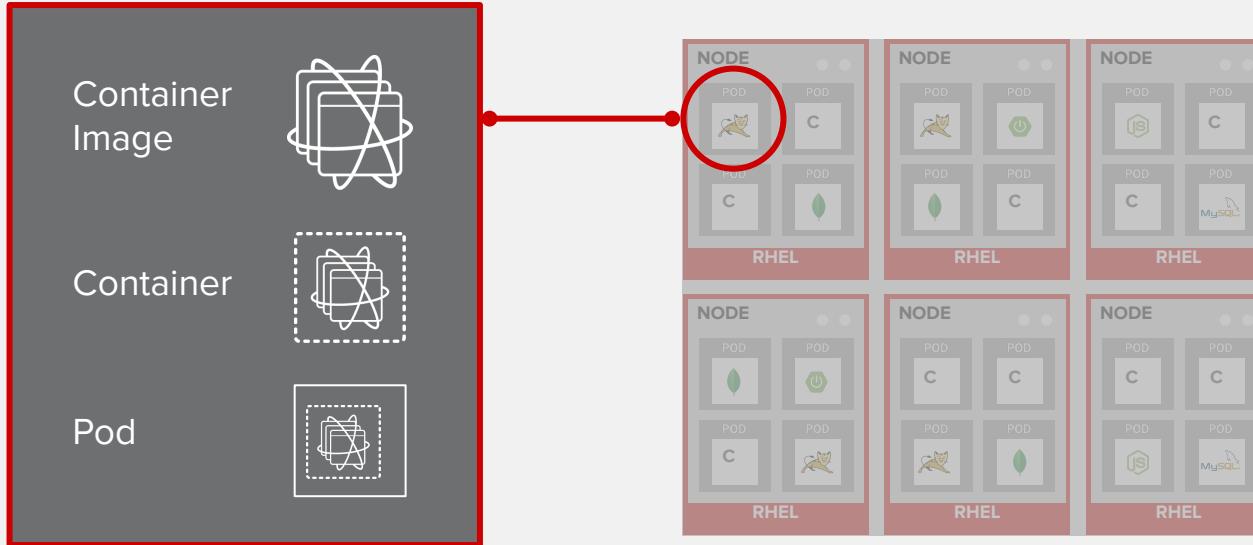
# Kubernetes Resource Definitions

What types of  
workloads can you  
deploy on top of a  
Container Platform...

# projects isolate apps across environments, teams, groups and departments



# APPS RUN IN CONTAINERS

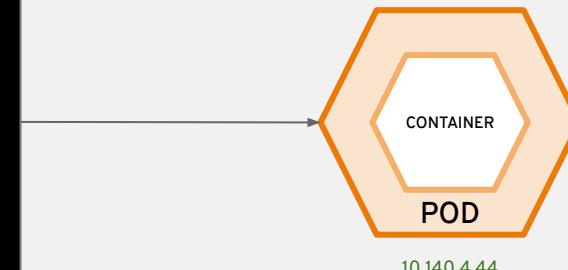


```
> oc get pods --all-namespaces  
> oc -n <namespace> describe pod <name>
```

# It all starts with an image

```
kind: ImageStream
metadata:
  name: lab-getting-started-session
spec:
  tags:
    - name: latest
  from:
    kind: DockerImage
    name: docker.io/httpd:latest
```

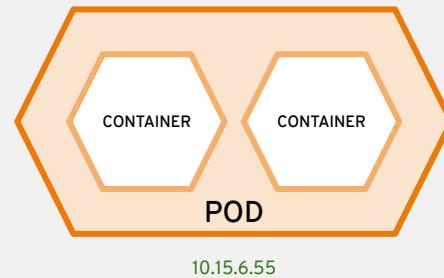
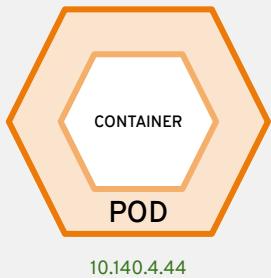
```
kind: Pod
metadata:
  name: mydb
spec:
  spec:
    containers:
      - name: backend
        image: mysql
        ports:
          - containerPort: 3306
        volumeMount:
          - name: data
            mount: /var/lib/mysql
        volumes:
          - name: data
            claim:
              requests:
                storage: 100Gi
```



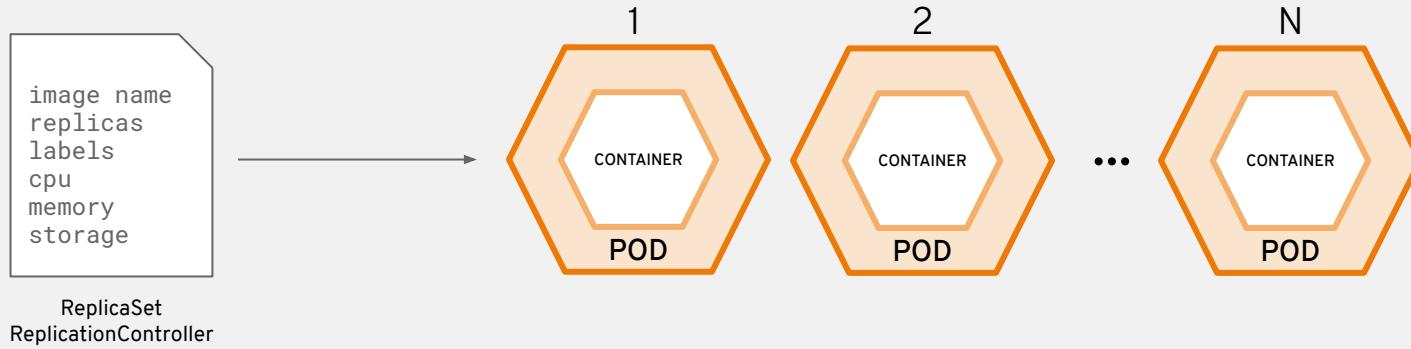
10.140.4.44

```
> oc -n openshift get imagestreams
> oc -n openshift describe imagestream httpd
> oc import-image
```

containers are wrapped in pods which are units of deployment and management



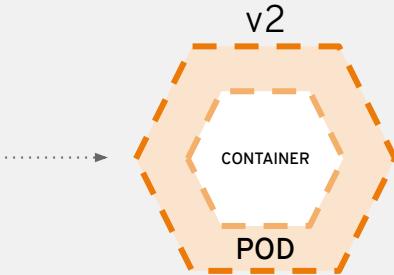
# ReplicationControllers & ReplicaSets ensure a specified number of pods are running at any given time



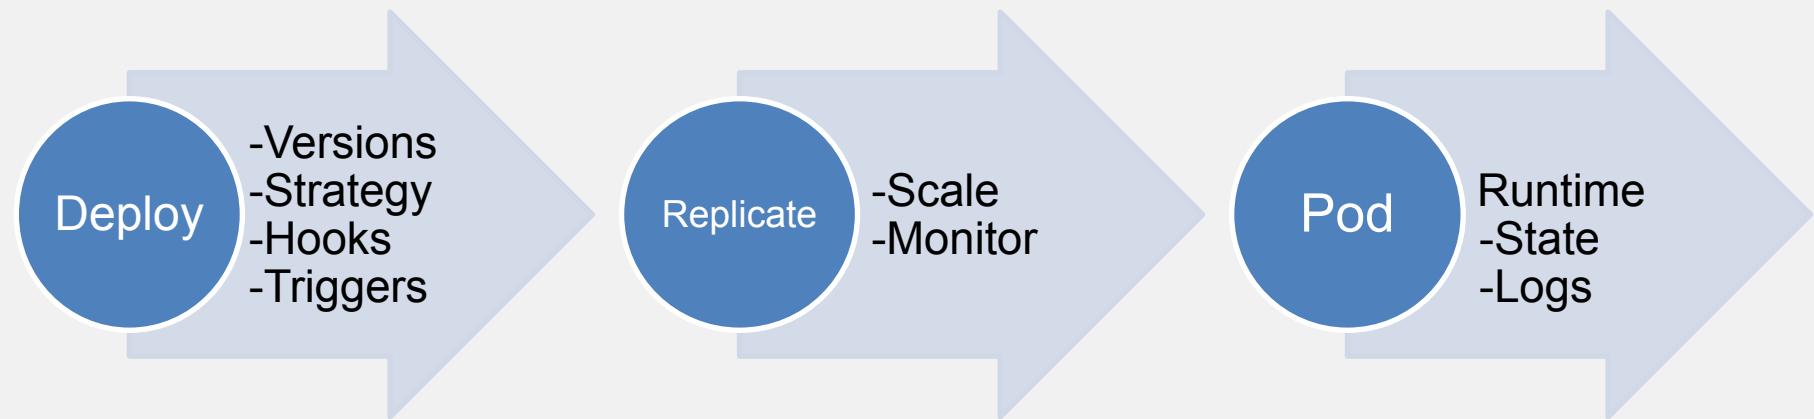
```
kind: "DeploymentConfig"  
metadata:  
  name: "myApp"  
spec:  
  replicas: 2  
  selector:  
    app: myapp  
  template:  
    metadata:  
      name: myapp  
      labels:  
        app: mine  
    spec:  
      containers:  
        - name: frontend  
          image: jboss-eap:latest  
          ports:  
            - containerPort: 80  
triggers:  
  - type: "ImageChange"  
    from:  
      kind: "Image"  
      name: "myapp:latest"
```

```
kind: "DeploymentConfig"  
metadata:  
  name: "myApp"  
spec:  
  replicas: 2  
  template:  
    spec:  
      containers:  
        - name: frontend  
        - name: flyway  
  strategy:  
    type: rolling  
  rollingParams:  
    pre:  
      execNewPod:  
        containerName: flyway  
        volumes: ['git']  
        command: "flyway do"  
    post:  
      tagImage:  
        containerName: frontend  
        to: "frontend:prod"  
triggers: ...
```

and  
ions define how to  
s of Pods



# Deployment Process



- MyJBossApp

```
> oc new-app httpd  
> oc scale -replicas=2 dc/frontend  
> oc rollout latest dc/frontend
```

- MyJBossApp-v1 (2x)
- MyJBossApp-v2 (4x)

```
> oc get pods -o wide -w  
> curl http://<pod-ip>:8080/
```

- MyJBossApp-v1-abcde
- MyJBossApp-v1-rando

```
> oc new-app httpd  
> oc scale -replicas=2 dc/frontend  
> oc rollout latest dc/frontend  
> oc rollback dc/frontend
```

```
> oc edit dc/frontend  
> oc describe pod/<pod-name>
```

```
> oc logs <pod-name>  
> oc rsh <pod-name> bash
```

```
> oc get pods -o wide -w  
> curl http://<pod-ip>:8080/
```

```
$ oc edit configmap scheduler-policy -n openshift-config
```

```
apiVersion: v1
```

```
data:
```

```
policy.cfg: |
```

```
{  
    "kind" : "Policy",  
    "apiVersion" : "v1",  
    "predicates" : [  
        {"name" : "MaxGCEPDVolumeCount"},  
        {"name" : "GeneralPredicates"},  
        {"name" : "MaxAzureDiskVolumeCount"},  
        {"name" : "MaxCSIVolumeCountPred"},  
        {"name" : "CheckVolumeBinding"},  
        {"name" : "MaxEBSVolumeCount"},  
        {"name" : "PodFitsResources"},  
        {"name" : "MatchInterPodAffinity"},  
        {"name" : "CheckNodeUnschedulable"},  
        {"name" : "NoDiskConflict"},  
        {"name" : "NoVolumeZoneConflict"},  
        {"name" : "MatchNodeSelector"},  
        {"name" : "HostName"},  
        {"name" : "PodToleratesNodeTaints"}  
    ],
```

# PLACEMENT BY POLICY

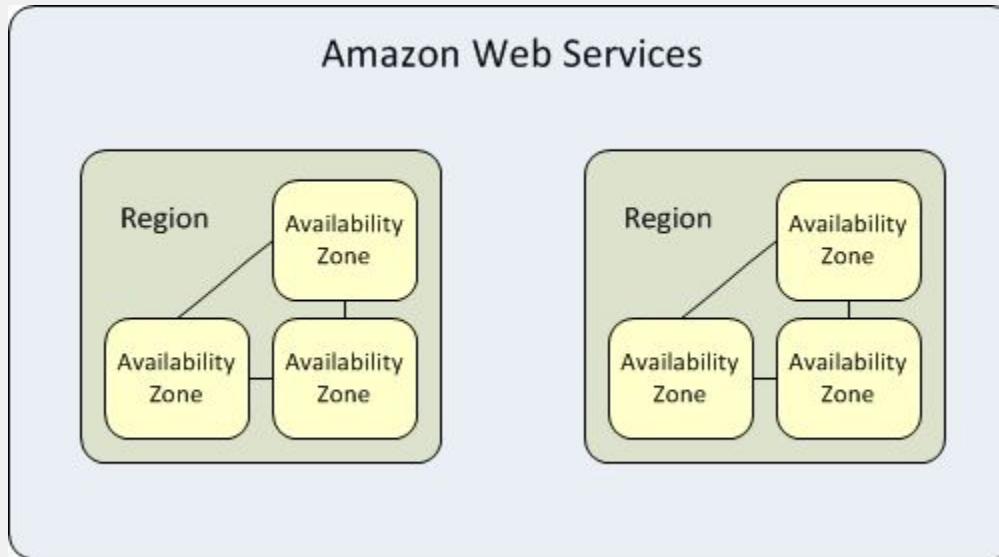
```
"priorities" : [
```

```
    {"name" : "LeastRequestedPriority", "weight"  
    {"name" : "BalancedResourceAllocation", "wei  
    {"name" : "ServiceSpreadingPriority", "wei  
    {"name" : "NodePreferAvoidPodsPriority", "wei  
    {"name" : "NodeAffinityPriority", "weight" :  
    {"name" : "TaintTolerationPriority", "weight"  
    {"name" : "ImageLocalityPriority", "weight" :  
    {"name" : "SelectorSpreadPriority", "weight"  
    {"name" : "InterPodAffinityPriority", "weight"  
    {"name" : "EqualPriority", "weight" : 1}  
]
```

```
}
```

# PLACEMENT BY POLICY

## nodeLabels: Zone + Region



Pods belonging to same “Service”:

MustHave:

Affinity:  
Region

ShouldHave:

Anti-Affinity:  
Zone (aka failure-domain)

```
oc get nodes --show-labels
```

```
> oc label node node01.apeldoorn.internal zone=A  
> oc label node node02.apeldoorn.internal zone=B  
> oc scale --replicas=2 dc/httpd
```

```
> oc -n userXX edit dc/httpd
```

```
spec:
```

```
template:
```

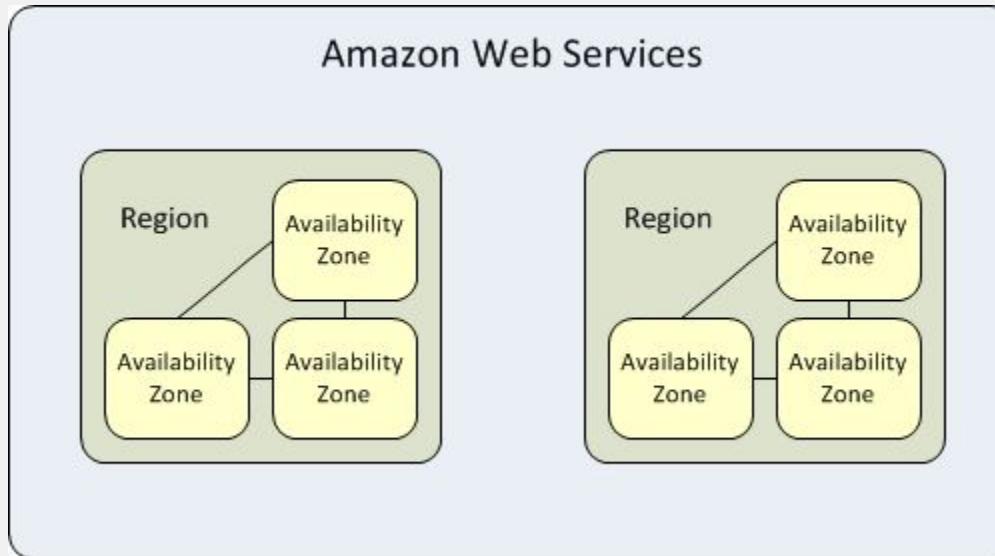
```
spec:
```

```
nodeSelector:
```

```
zone: A
```

# PLACEMENT BY POLICY

## Pod Affinity rules



Specify complex query

# PLACEMENT BY POLICY

## Resource Reservations

...

---

```
> ssh master1 cat /etc/origin/master/scheduler.json
```

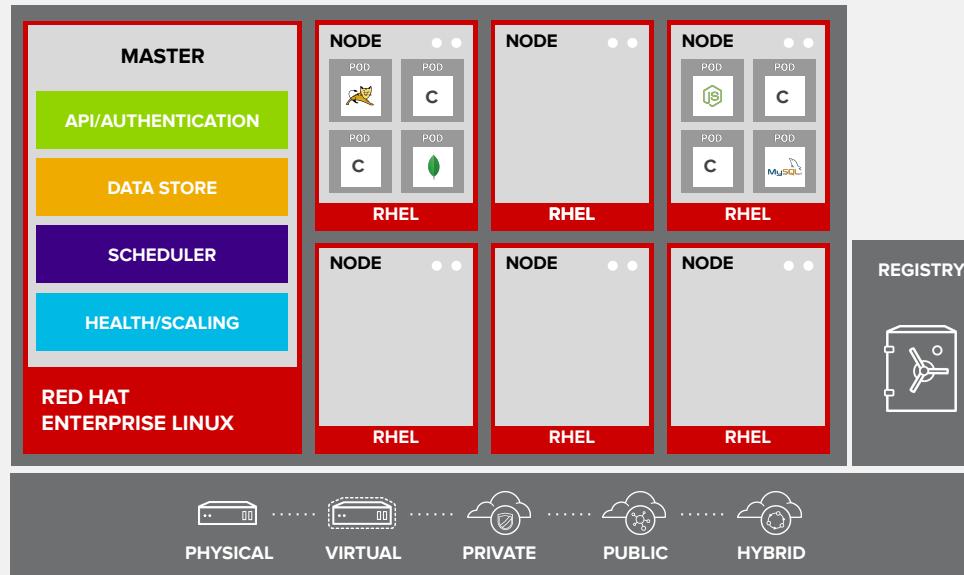
# Resource Limits

- DefaultLimits

---

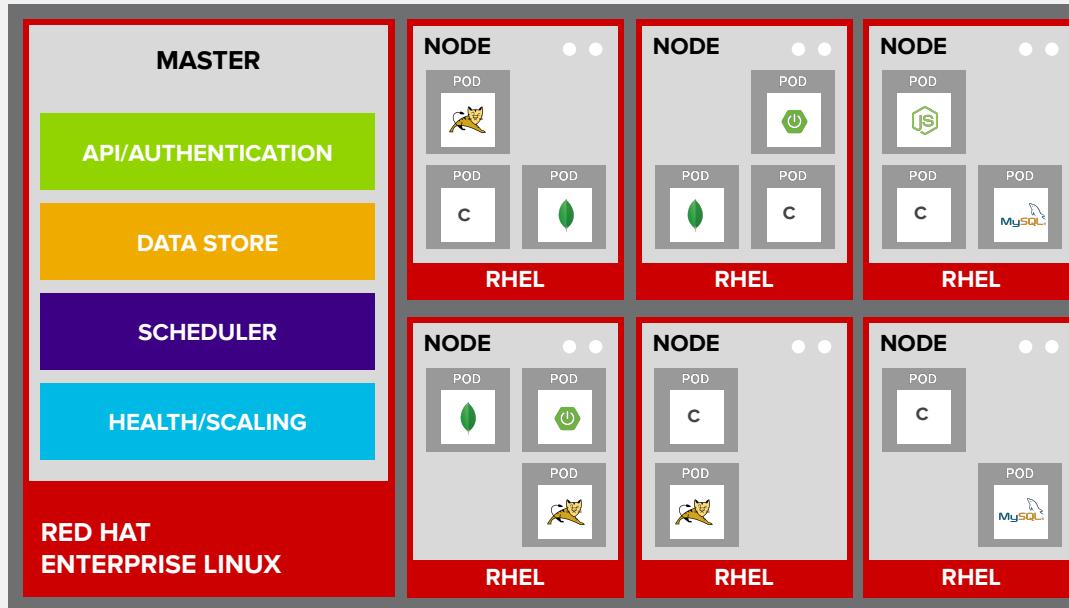
```
> ssh master1 cat /etc/origin/master/scheduler.json
```

# HEALTH AND SCALING



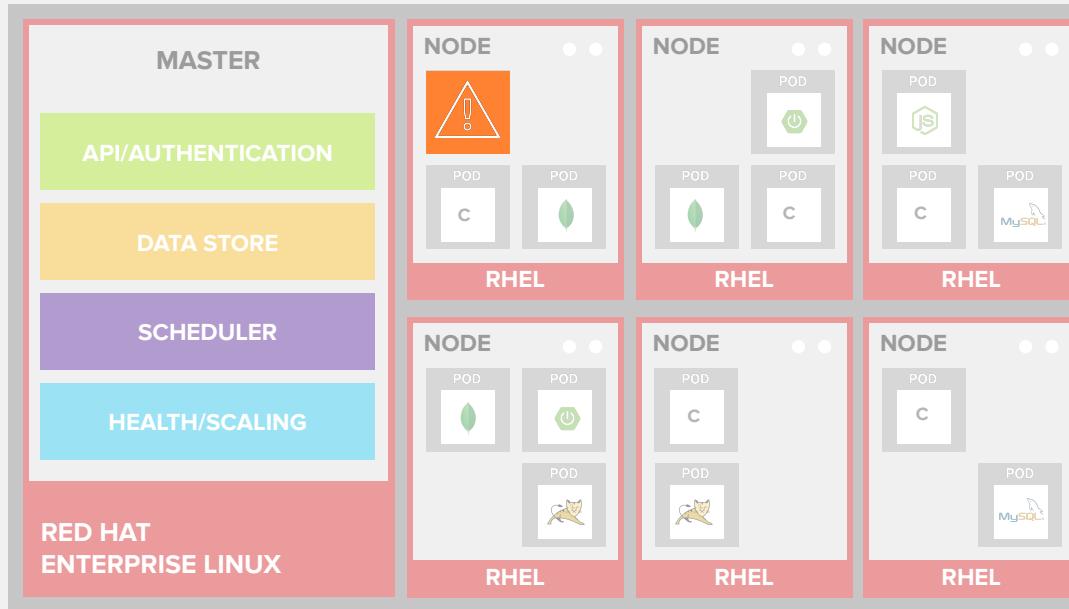
```
> oc describe dc/jenkins legrep "Liveness|Readiness"  
> oc scale --replicas=4 dc/jenkins; oc get pods --watch
```

# AUTO-HEALING FAILED CONTAINERS



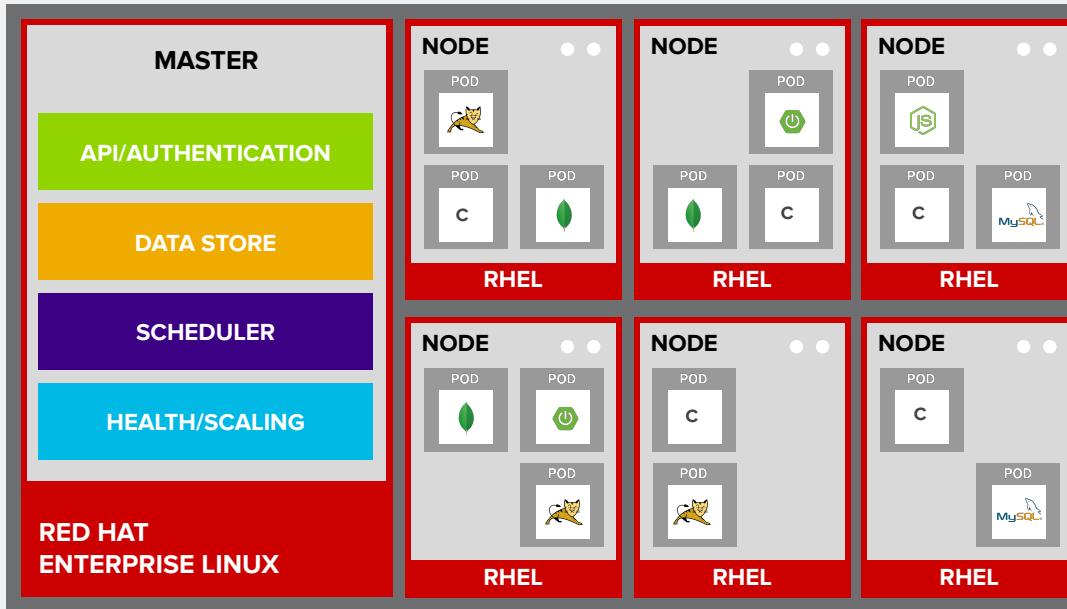
```
> oc get pods -o wide --watch &  
> oc delete pods --all
```

# AUTO-HEALING FAILED CONTAINERS



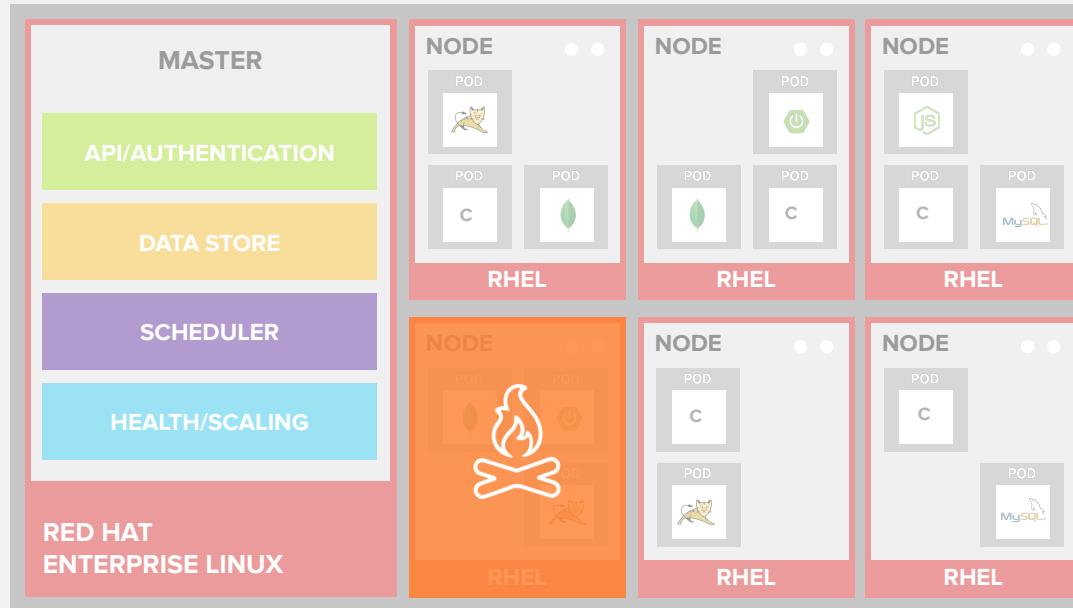
```
> oc get pods -o wide --watch &  
> oc delete pods --all
```

# AUTO-HEALING FAILED CONTAINERS



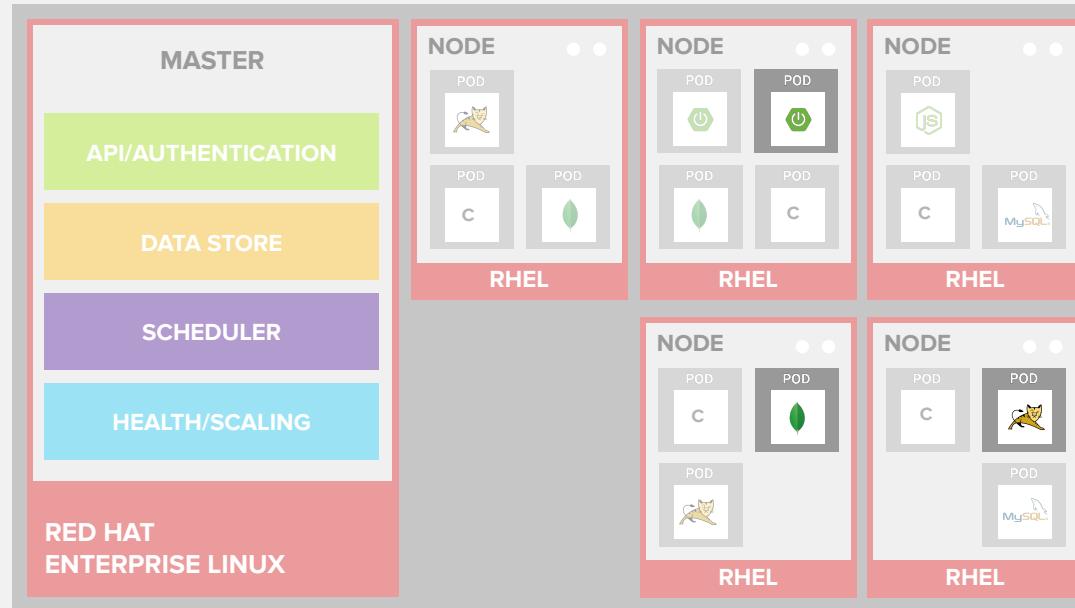
```
> oc get pods -o wide --watch &  
> oc delete pods --all
```

# AUTO-HEALING FAILED NODES



```
> oc get pods -o wide --watch &  
> oc delete pods --all
```

# AUTO-HEALING FAILED NODES



```
> oc get pods -o wide --watch &  
> oc delete pods --all
```

# Health Checks

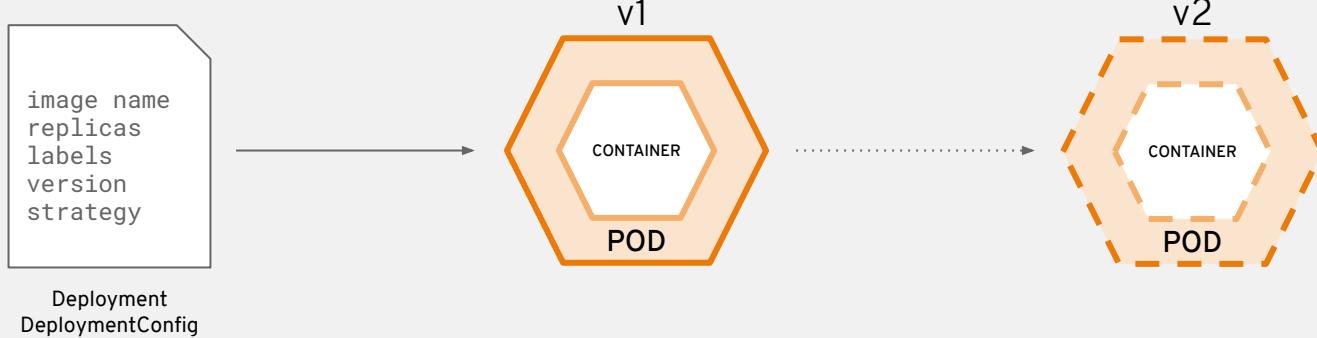
- LivenessProbe
- ReadinessProbe

HTTP  
TCP  
Shell

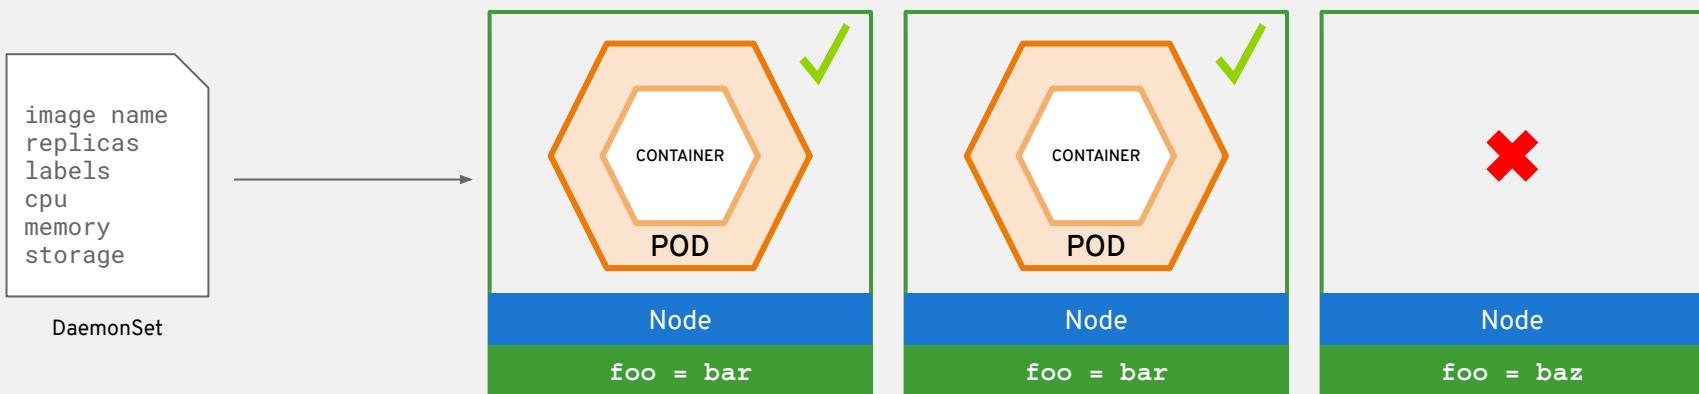
---

```
> oc set probe dc/httpd --readiness --get-url=http://127.0.0.1:8080/index.html --initial-delay-seconds=5
> oc set probe dc/httpd --liveness --open-tcp=8080
```

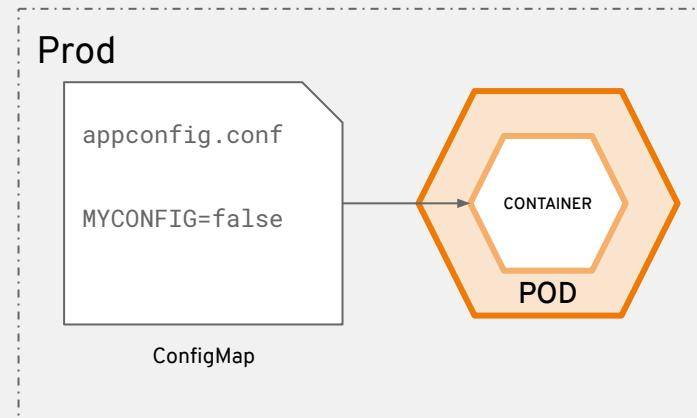
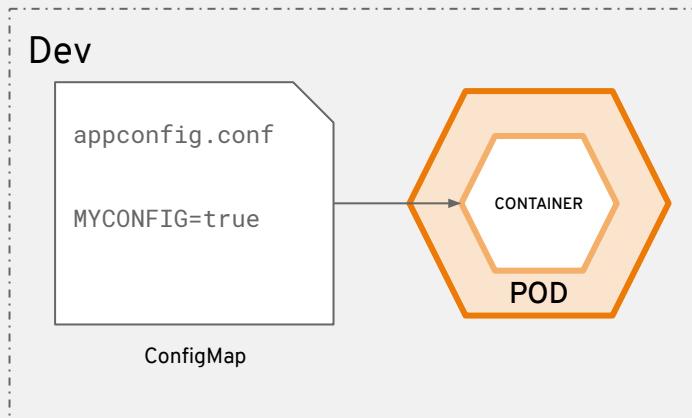
# StatefulSets are special deployments for Stateful workloads



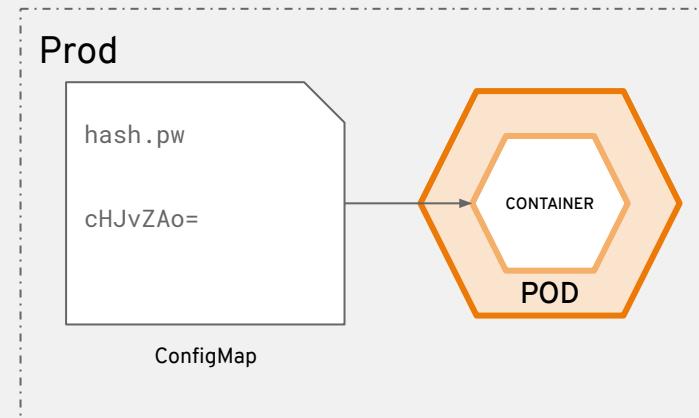
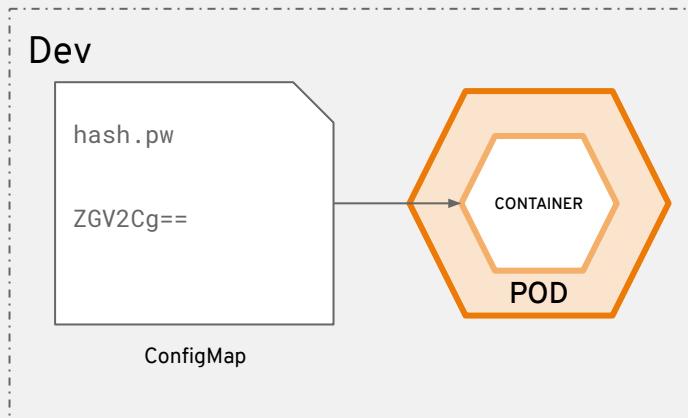
a daemonset ensures that all  
(or some) nodes run a copy of a pod



configmaps allow you to decouple configuration artifacts from image content

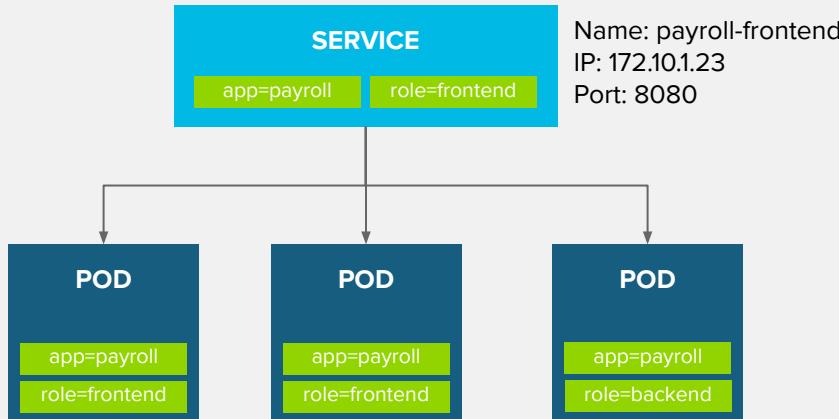


secrets provide a mechanism to hold sensitive information such as passwords



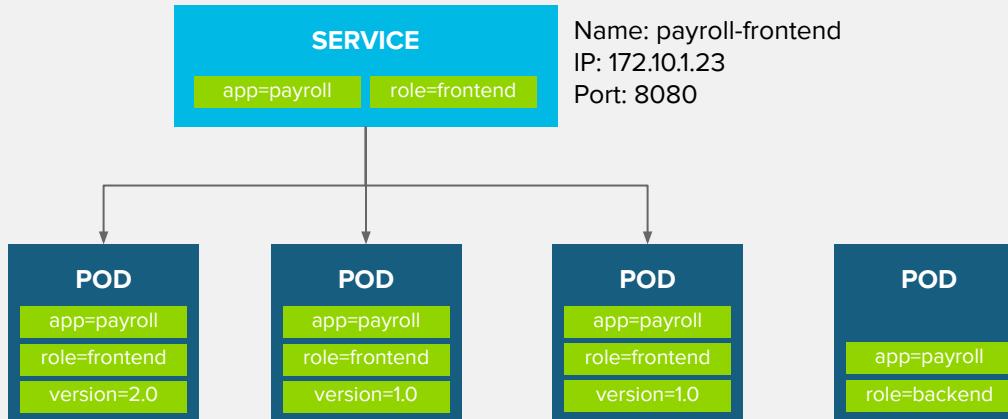
# Networking

# BUILT-IN SERVICE DISCOVERY & INTERNAL LOAD-BALANCING



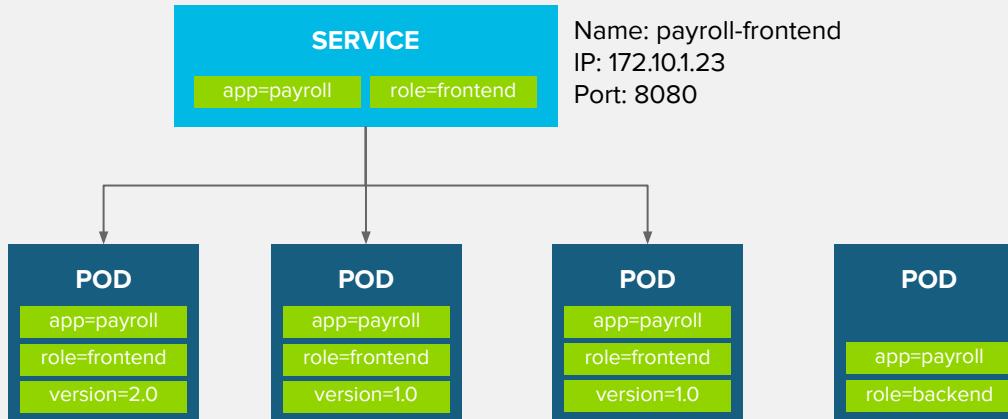
```
> oc describe svc/jenkins legrep "Selector|Endpoints"  
> oc get pods -l name=jenkins
```

# BUILT-IN SERVICE DISCOVERY & INTERNAL LOAD-BALANCING



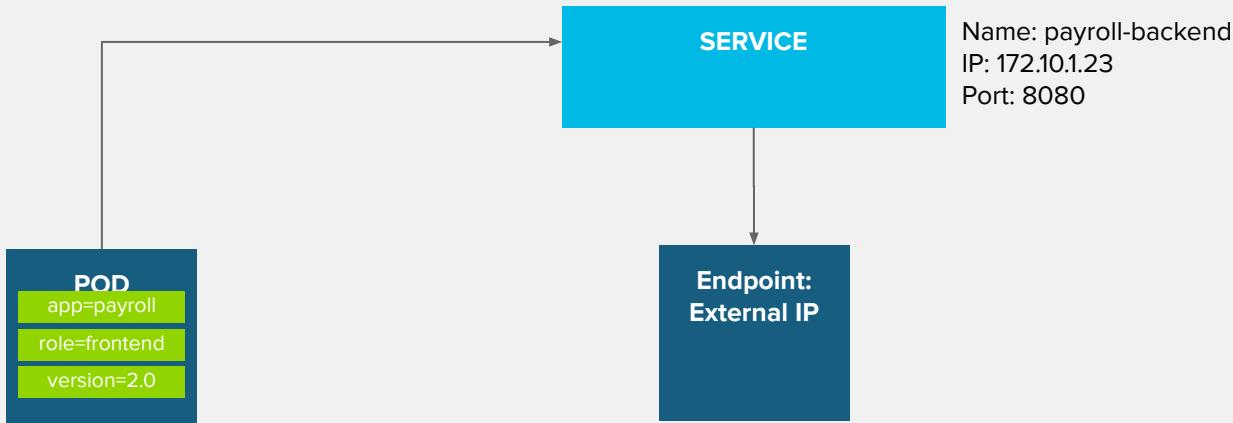
```
> oc describe svc/jenkins legrep "Selector|Endpoints"  
> oc get pods -l name=jenkins
```

# BUILT-IN SERVICE DISCOVERY & INTERNAL LOAD-BALANCING



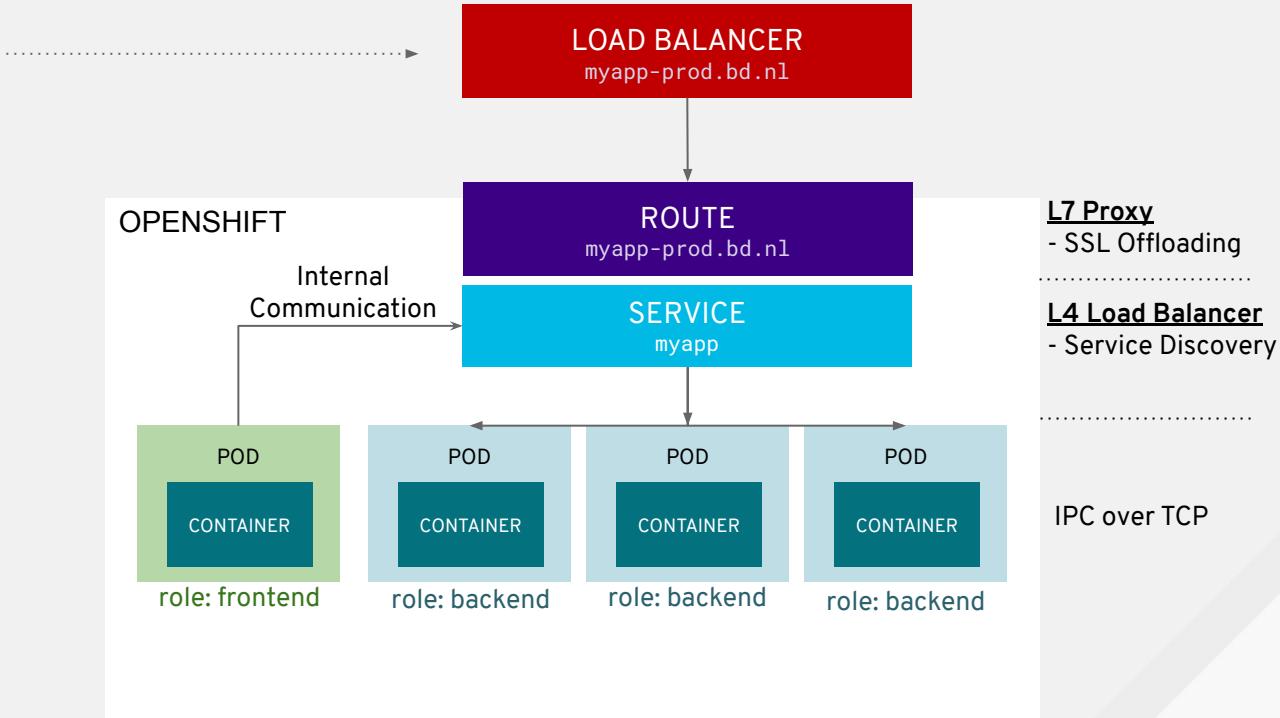
```
> oc describe svc/jenkins legrep "Selector|Endpoints"  
> oc get pods -l name=jenkins
```

# EXTERNAL SERVICES



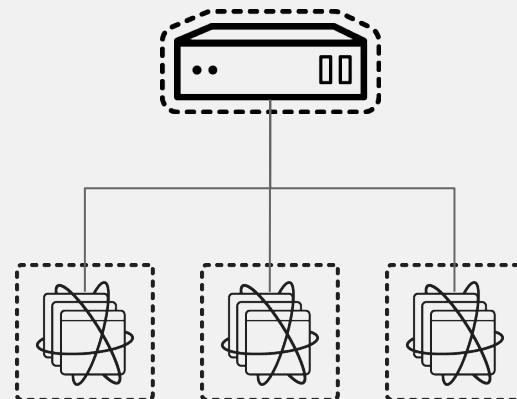
```
> oc describe svc/jenkins legrep "Selector|Endpoints"  
> oc get pods -l name=jenkins
```

# ROUTING TRAFFIC



# ROUTING AND EXTERNAL LOAD-BALANCING

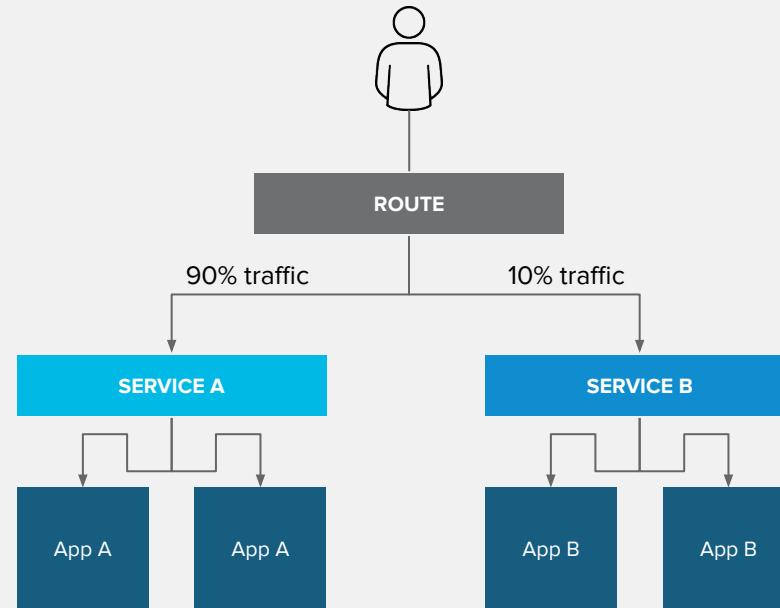
- Pluggable routing architecture
  - HAProxy Router
  - F5 Router
- Multiple-routers with traffic sharding
- Router supported protocols
  - HTTP/HTTPS
  - WebSockets
  - TLS with SNI
- Non-standard ports via cloud load-balancers, ExternalIP, and NodePort



```
> oc expose svc/jenkins --hostname jenkins-user00.openshift.eu  
> oc get route
```

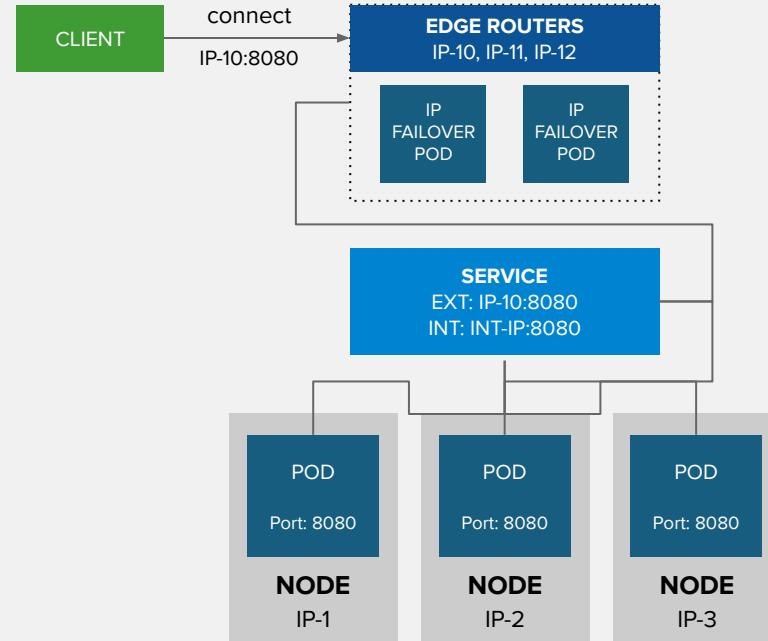
# ROUTE SPLIT TRAFFIC

Split Traffic Between  
Multiple Services For A/B  
Testing, Blue/Green and  
Canary Deployments



# EXTERNAL TRAFFIC WITH EXTERNALIP

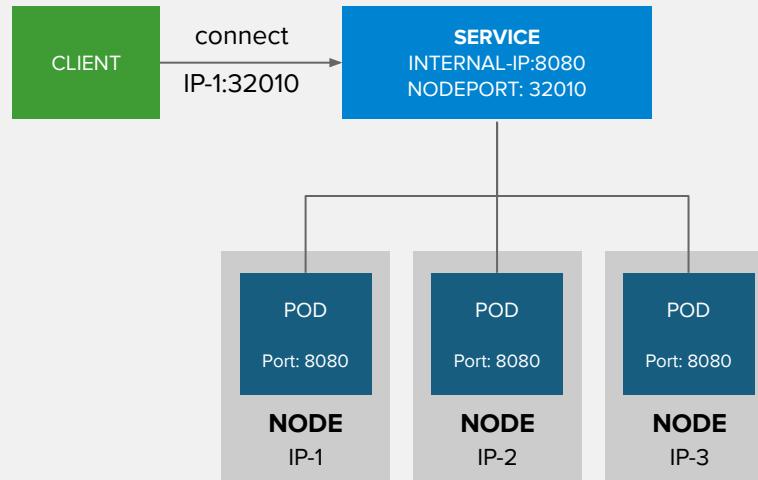
- Route external traffic to a service on any TCP/UDP port
- Available on non-cloud clusters
- External IP automatically assigned from a pre-defined pool of external IPs
- IP failover pods provide high availability for the pool of external IPs



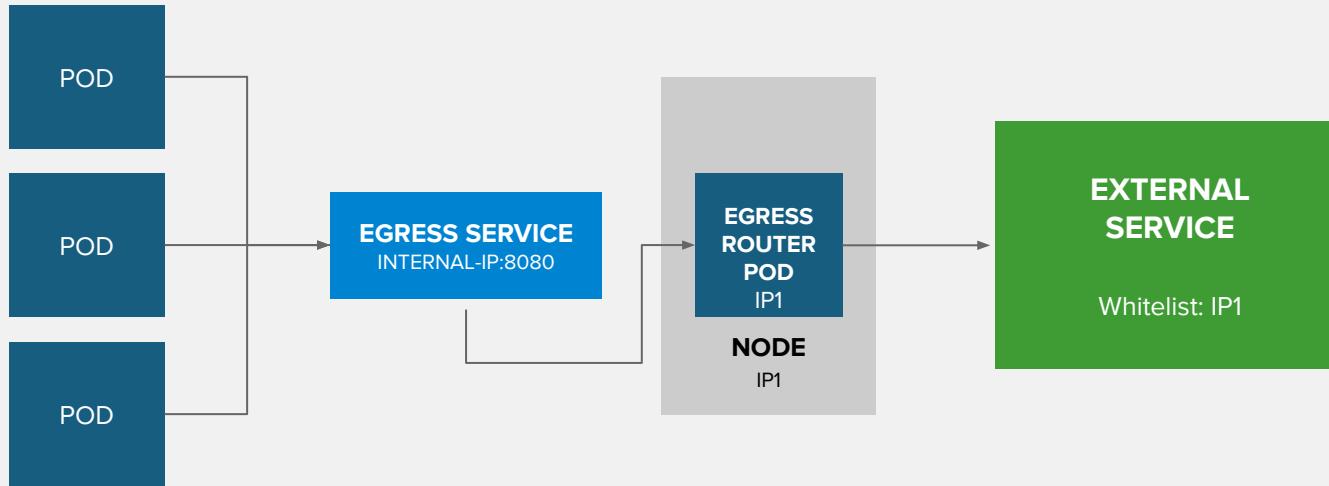
```
> oc expose dc/jenkins --external-ip=172.29.253.100 --name=jenkins-externalip  
> oc get svc
```

# EXTERNAL TRAFFIC WITH NODEPORT

- NodePort exposes a unique port on all the nodes in the cluster
- Ports in 30K-60K range which usually differs from the service
- Traffic received on any node redirects to a node with the running service
- Firewall rules must allow traffic to all nodes on the specific port



# CONTROL SOURCE IP WITH EGRESS ROUTER

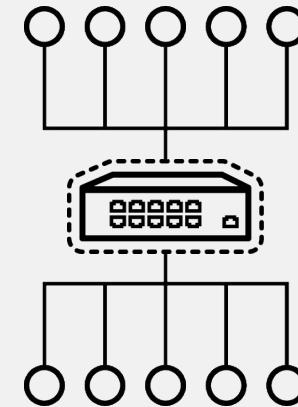


# OpenShift Networking

Features, mechanisms  
and processes for  
container and platform  
isolation

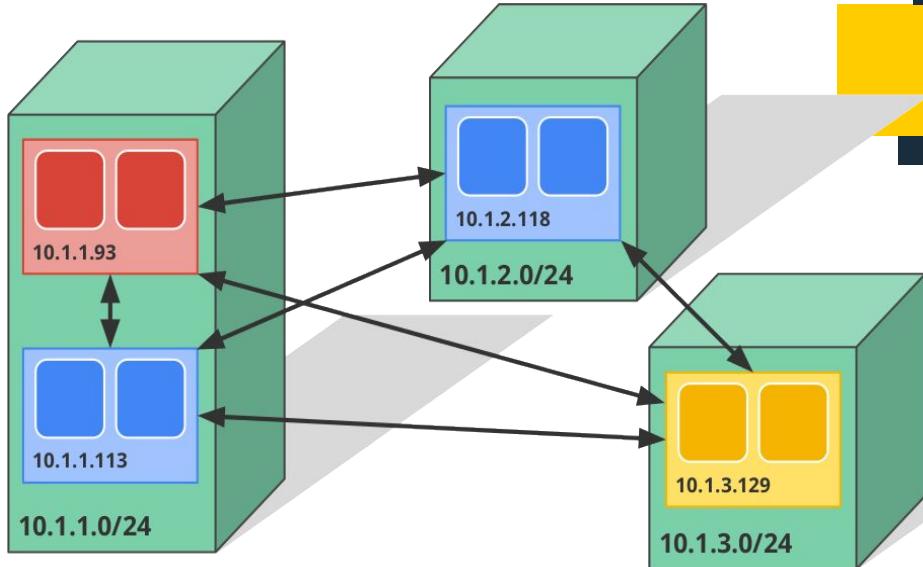
# OPENSHIFT NETWORKING

- Built-in internal DNS to reach services by name
- Split DNS is supported via SkyDNS
  - Master answers DNS queries for internal services
  - Other nameservers serve the rest of the queries
- Software Defined Networking (SDN) for a unified cluster network to enable pod-to-pod communication
- OpenShift follows Kubernetes network plug-in model
- Supported plug-ins
  - OpenShift SDN (Open vSwitch or Flannel)
  - Nuage SDN (Virtualized Services Platform)



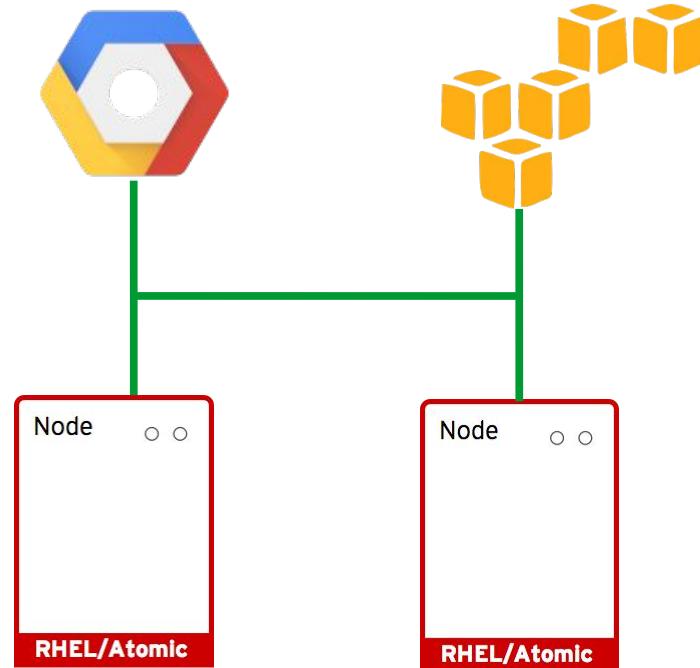
# Overlay Networking

- Each Host = /24 IP-Range
- Each POD = 1 IP
  - Containers in Pod = localhost
- Egress traffic = NAT Node's IP
- Network Plugins:
  - Flannel, Weave, Nuage
  - OpenShift-SDN (OpenVSwitch)
  - Neutron, GCE, AWS



# Underlay “infra” Network

- Programmable Infra
- Network Plugins (CNI):
  - GCE
  - AWS
  - OpenStack
  - Nuage
- Egress traffic = NAT Node IP



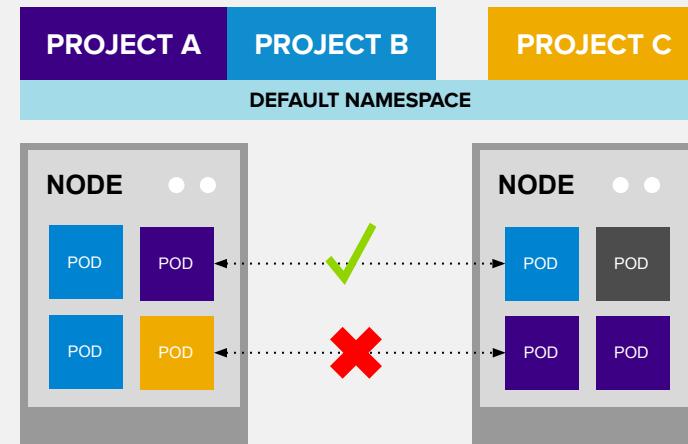
# OPENSHIFT SDN

## FLAT NETWORK

- All pods can communicate with each other across projects

## MULTI-TENANT NETWORK

- Project-level network isolation
- Granular policies for network traffic
- Multicast support
- Egress network policies



```
> oc adm pod-network --to=user01 user00  
> oc get netnamespaces lsort -k2 -n
```

# Persistent Storage

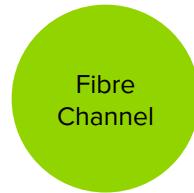
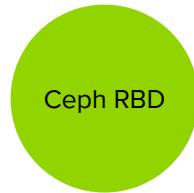
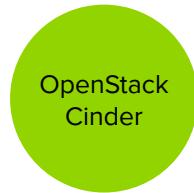
Connecting real-world storage to your containers to enable stateful applications

## A broad spectrum of static and dynamic storage endpoints

NFS	OpenStack Cinder	iSCSI	Azure Disk	AWS EBS	FlexVolume
GlusterFS	Ceph RBD	Fiber Channel	Azure File	GCE Persistent Disk	VMWare vSphere VMDK
NetApp Trident*		Container Storage Interface (CSI)**			

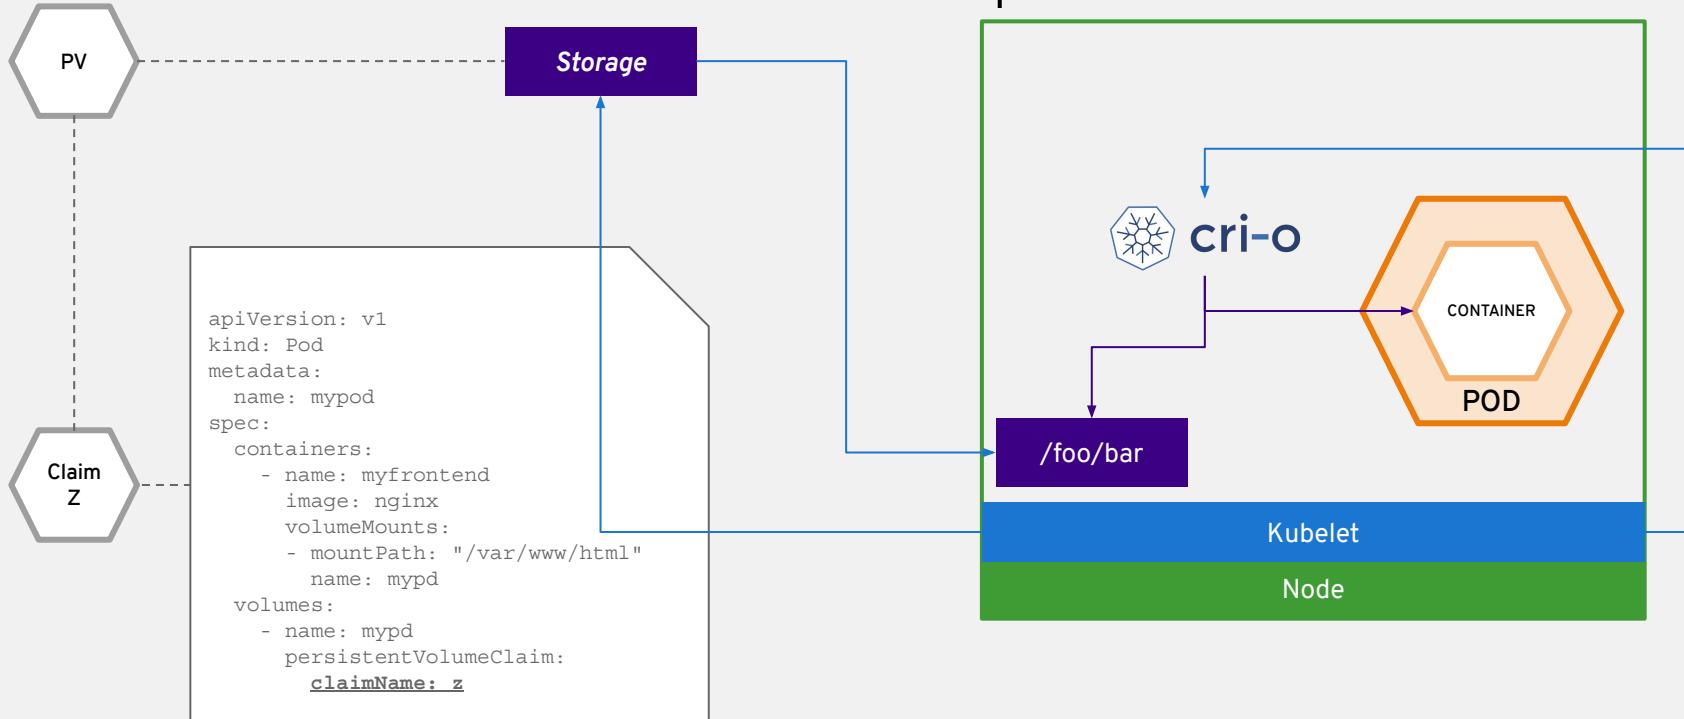
# PERSISTENT STORAGE

- Persistent Volume
  - Tied to a piece of network storage
  - Provisioned by an administrator (static or dynamically)
  - Allows admins to describe storage and users to request storage

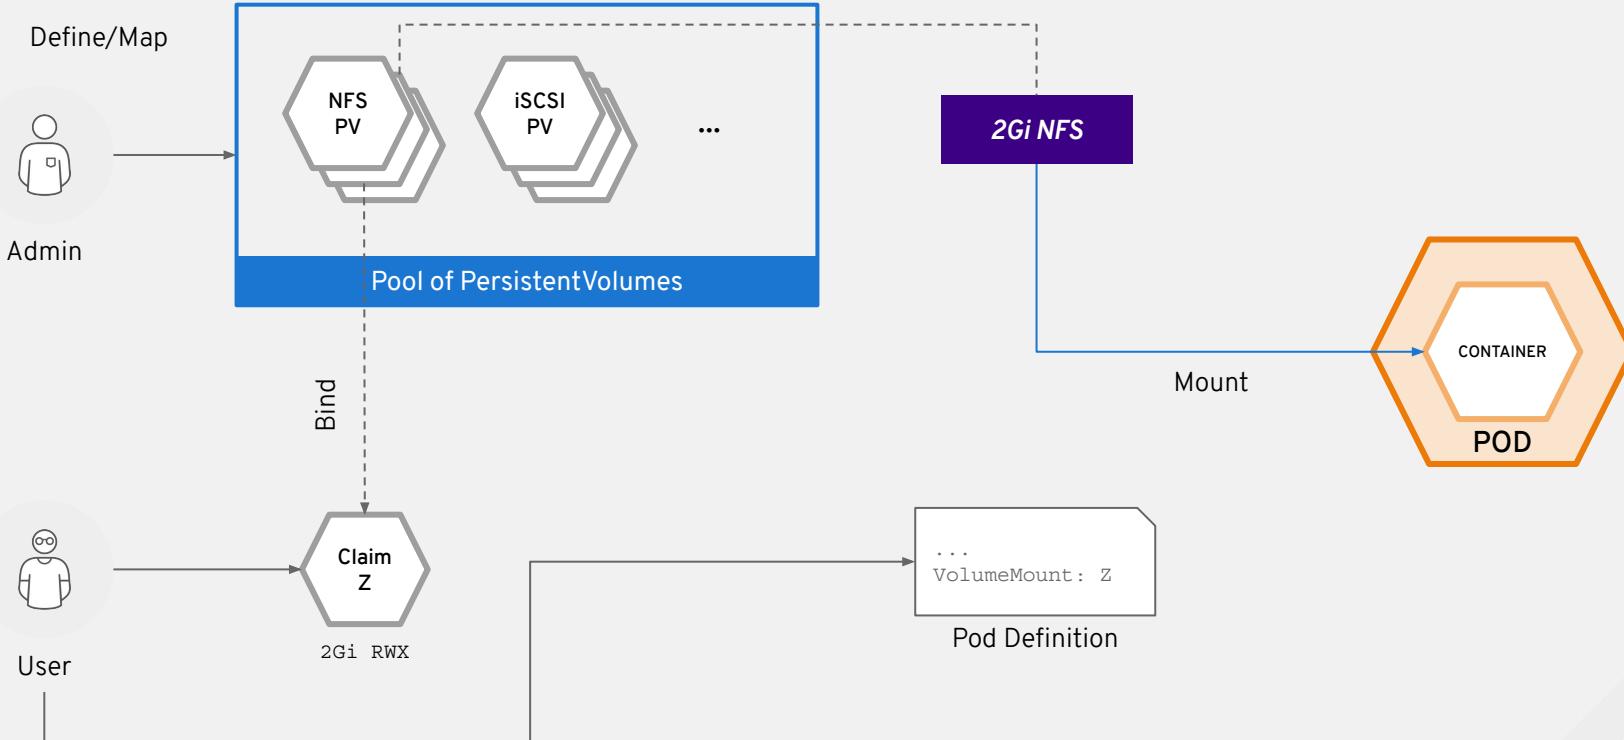


> oc get storageclass

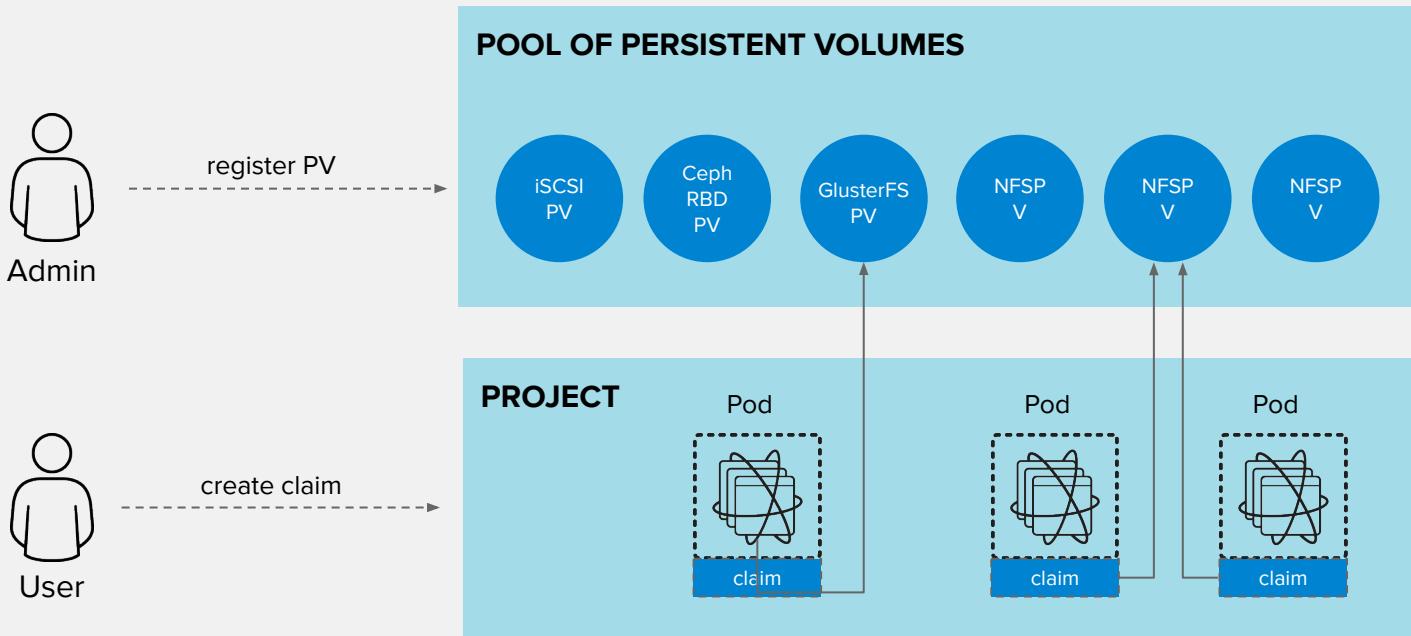
# PV Consumption



# Static Storage Provisioning

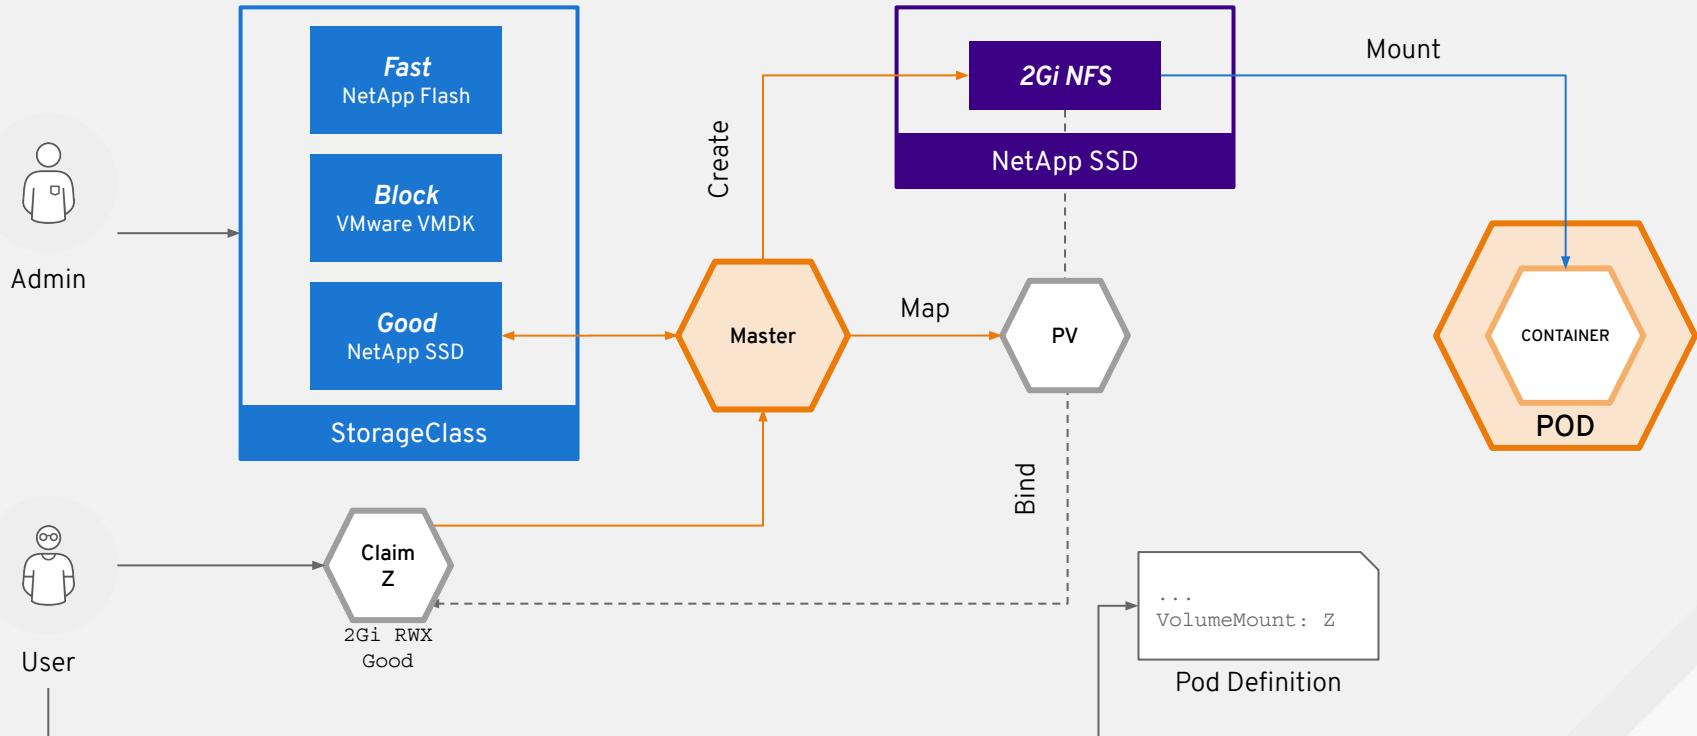


# PERSISTENT STORAGE

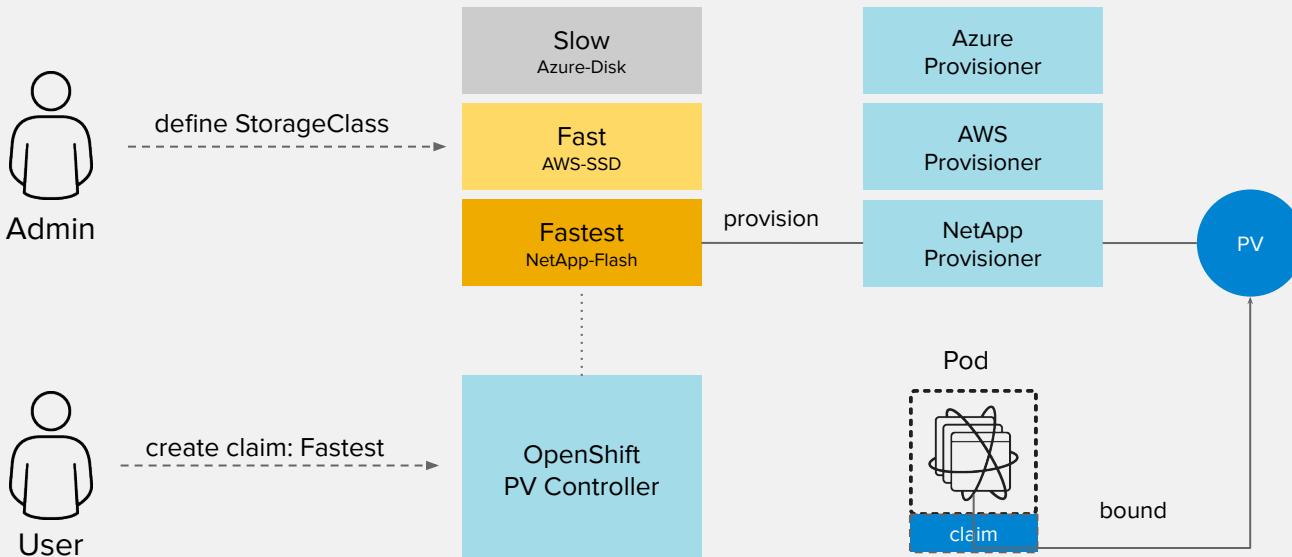


```
> oc get pv  
> oc get pvc --all-namespaces
```

# Dynamic Storage Provisioning



# DYNAMIC VOLUME PROVISIONING



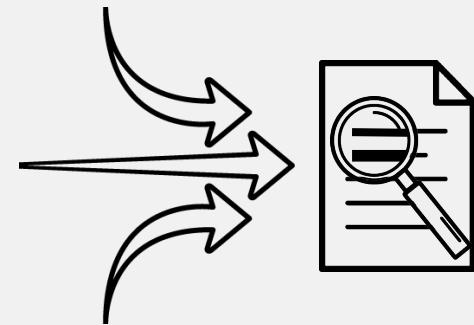
```
> oc get storageclass  
> oc -n glusterfs get all
```

# OpenShift Logging

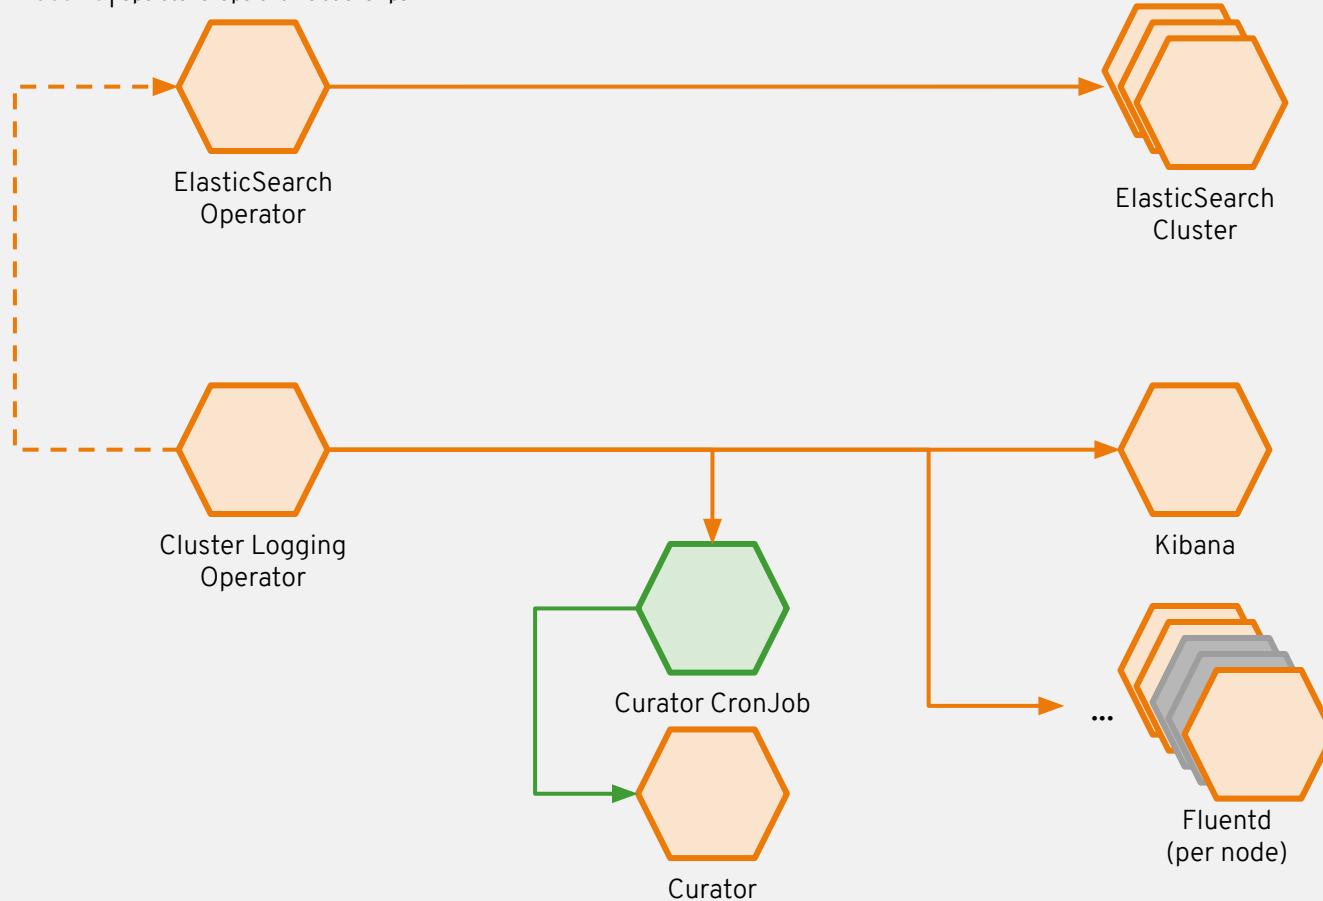
An integrated solution  
for exploring and  
corroborating  
application logs

# CENTRAL LOG MANAGEMENT WITH EFK

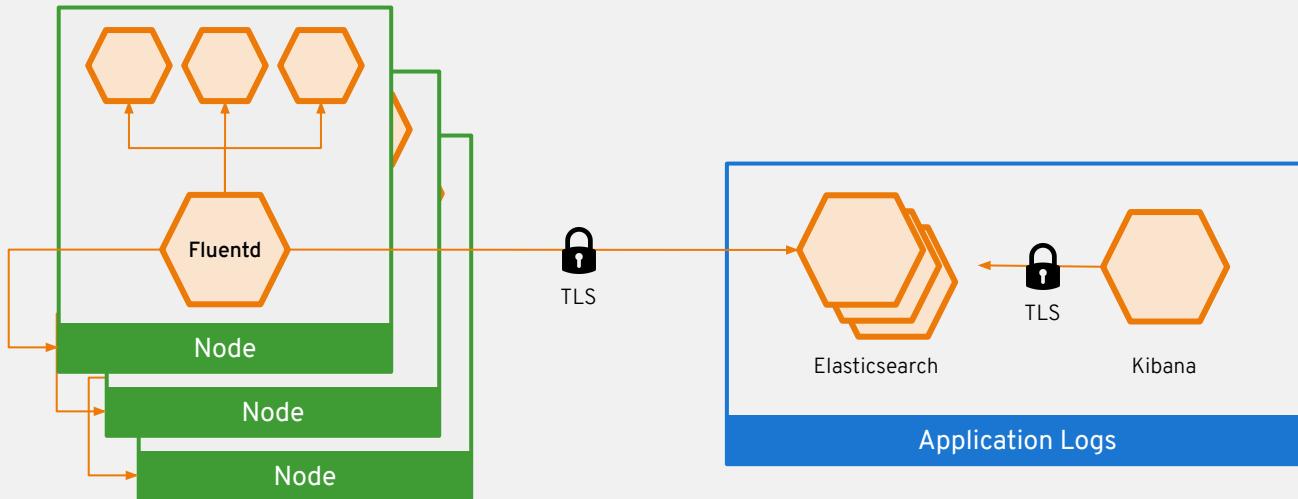
- EFK stack to aggregate logs for hosts and applications
  - **Elasticsearch:** an object store to store all logs
  - **Fluentd:** gathers logs and sends to Elasticsearch.
  - **Kibana:** A web UI for Elasticsearch.
- Access control
  - Cluster administrators can view all logs
  - Users can only view logs for their projects
- Ability to send logs elsewhere
  - External elasticsearch, Splunk, etc



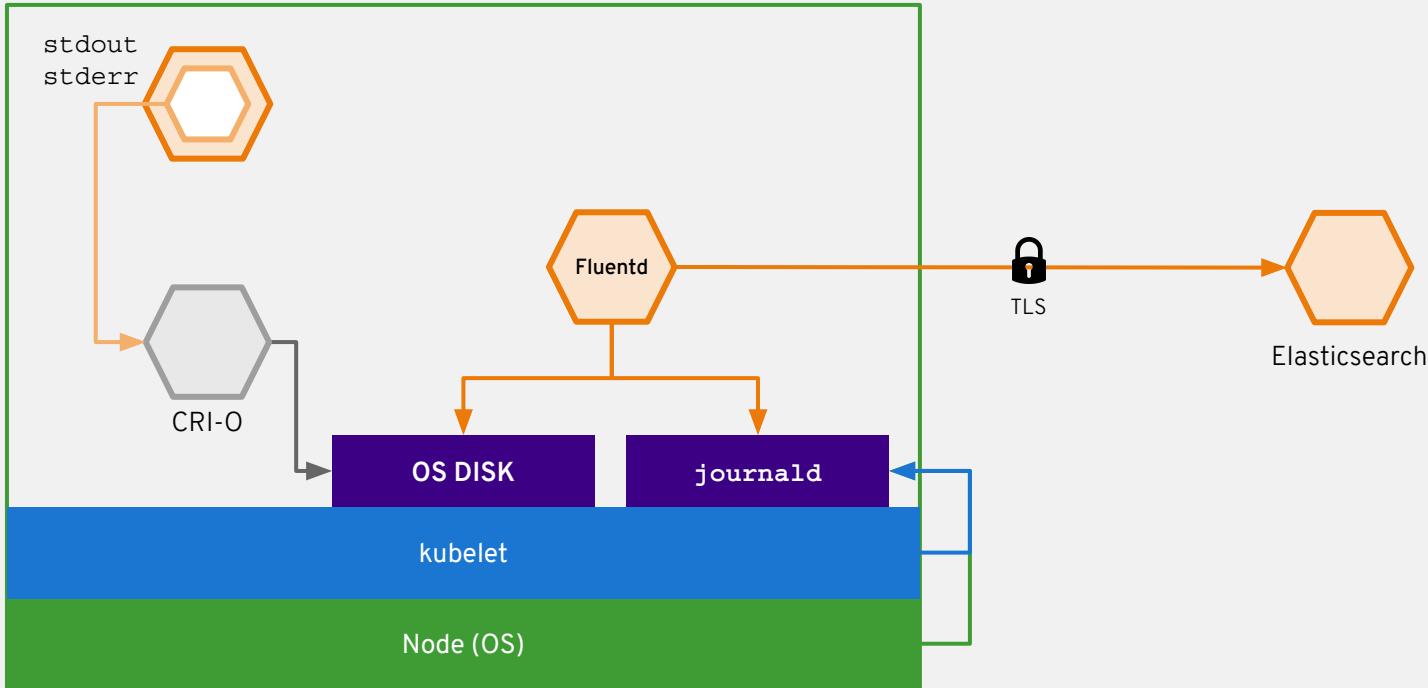
> oc -n logging get all  
> <http://kibana.openshift.eu>



# Log data flow in OpenShift



# Log data flow in OpenShift



# OpenShift Monitoring

An integrated cluster  
monitoring and alerting  
stack

# OpenShift Cluster Monitoring



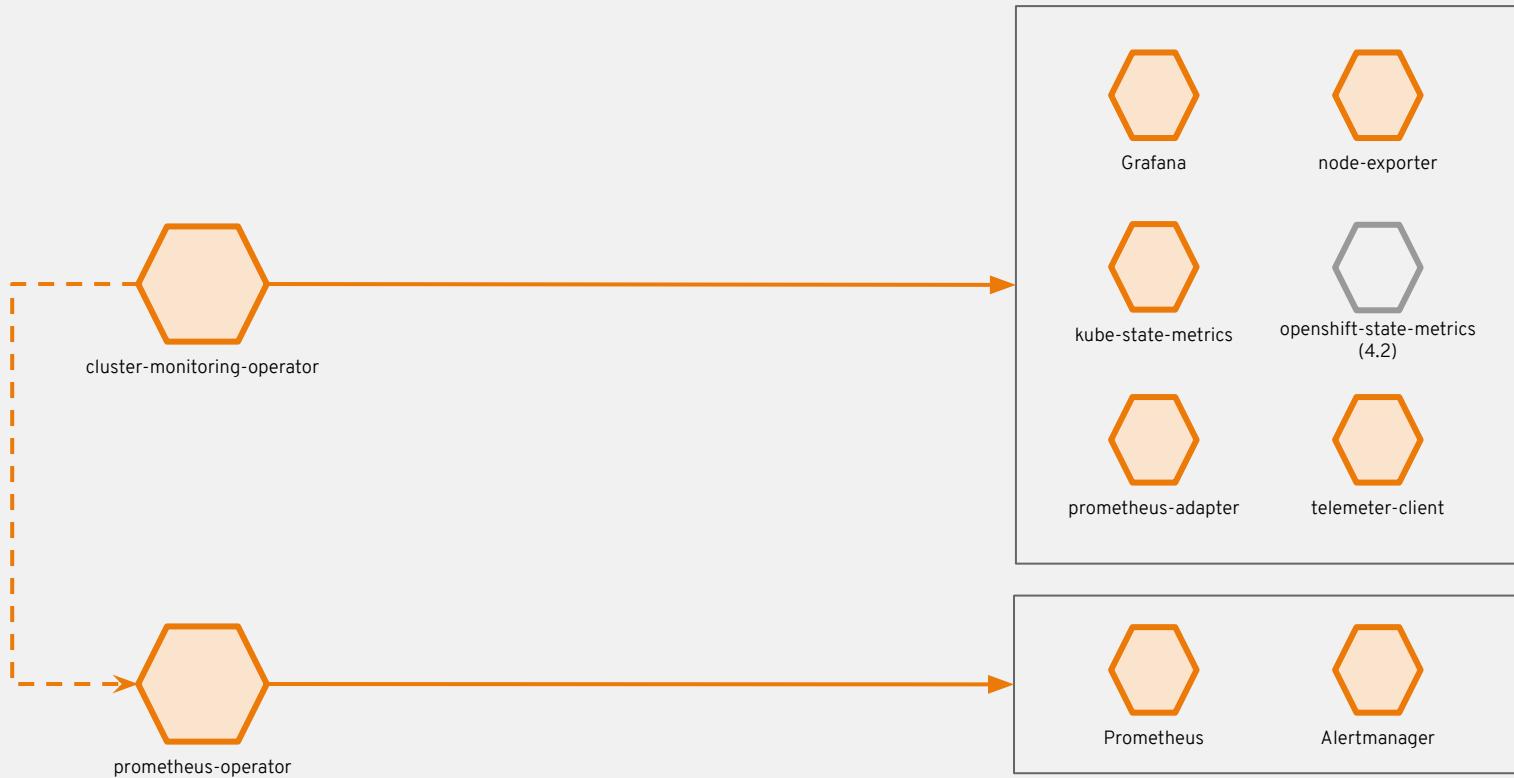
**Metrics collection and storage**  
via Prometheus, an  
open-source monitoring system  
time series database.

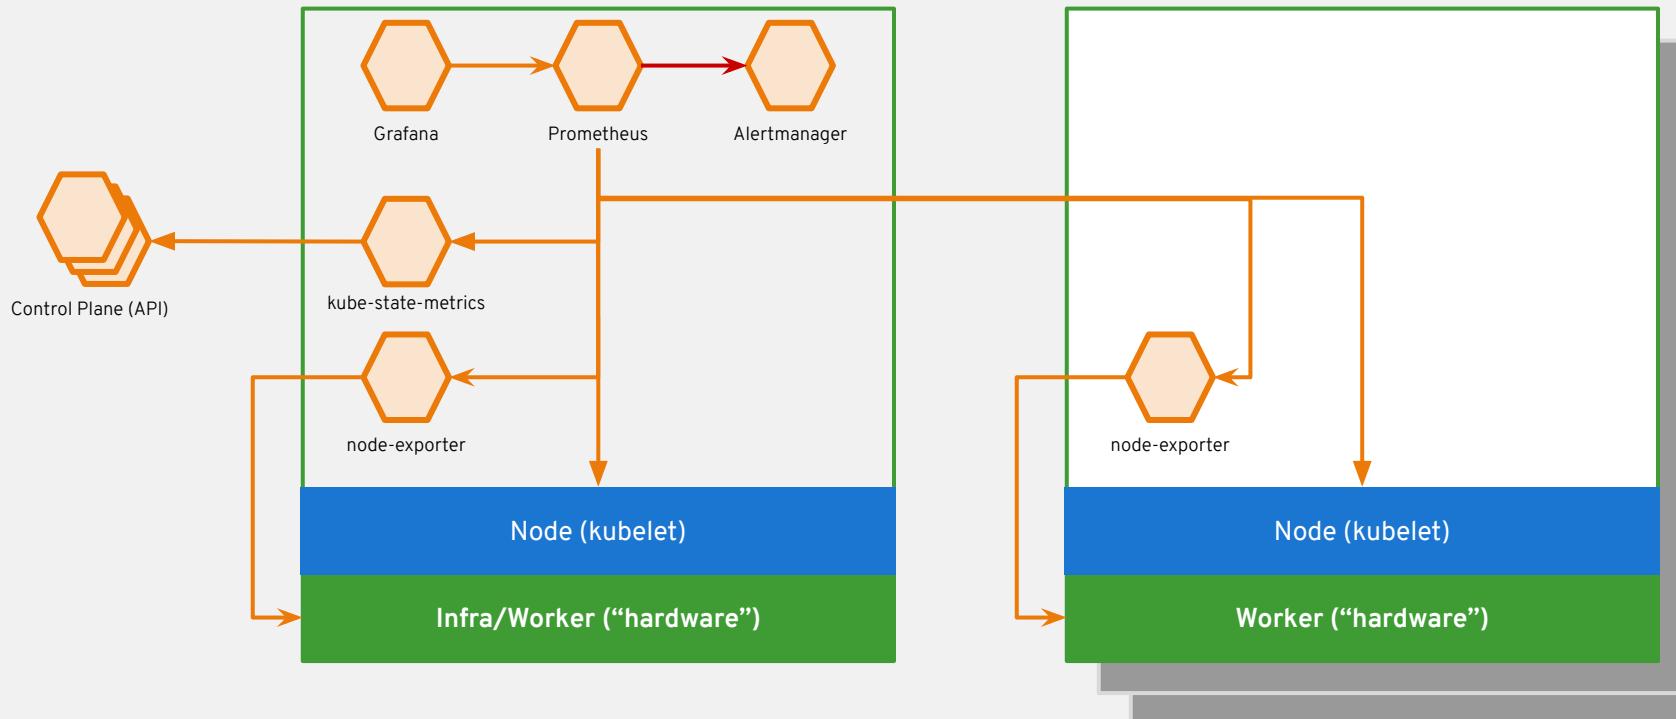


**Alerting/notification** via  
Prometheus' Alertmanager, an  
open-source tool that handles  
alerts send by Prometheus.



**Metrics visualization** via  
Grafana, the leading metrics  
visualization technology.





# **End of Day #1**

# AGENDA

## Day 1

09:00 Overview  
09:30 Architektur  
10:30 Tech. deep dive  
12:00 LUNCH  
13:00 Networking  
14:00 Logging / Metrics / Security  
14:30 Storage  
16:00 Operational Management  
16:45 Reference Architectures

## Day 2

09:00 Working with containers  
11:30 CI / CD Part 1  
12:30 LUNCH  
13:30 CI / CD Part 2  
15:00 Developer Workflow  
16:00 Application Services  
16:30 Q & A / next steps